

Advanced Programming in the UNIX Environment

Week 06, Segment 4: The Environment

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

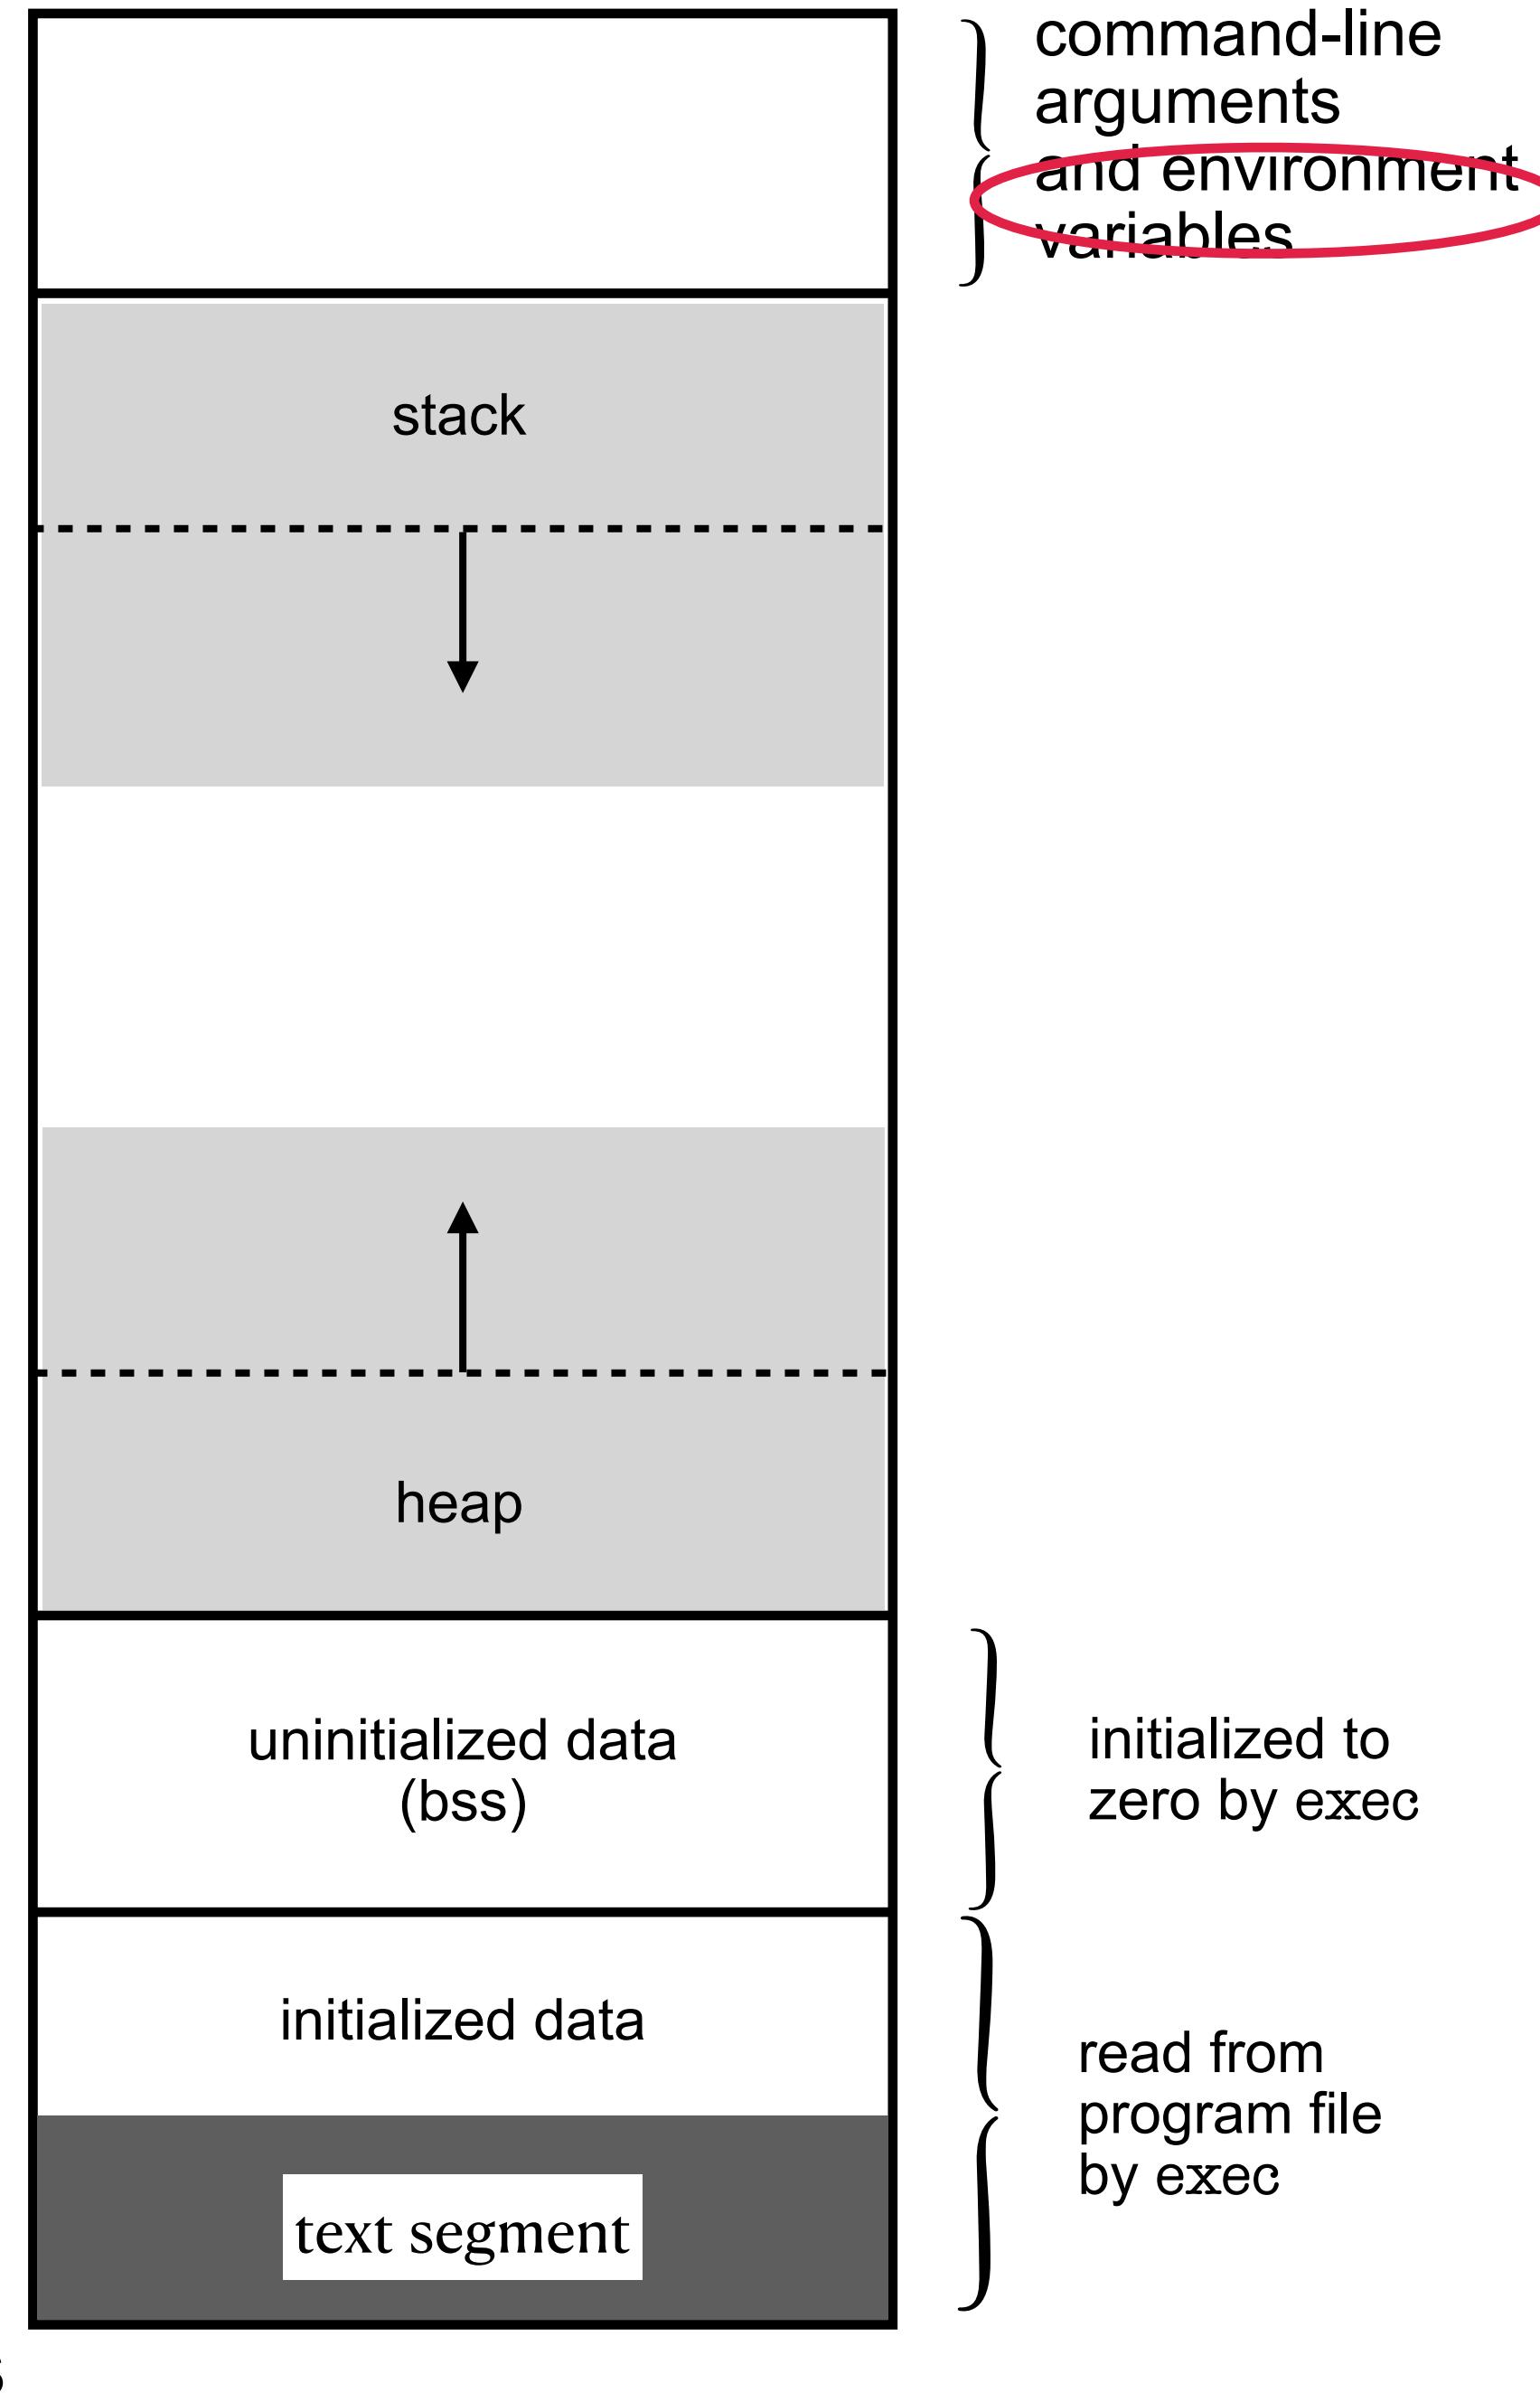
`https://stevens.netmeister.org/631/`



```
[apue$ env
ENV=/home/jschauma/.shrc
PS1=apue$
PWD=/home/jschauma/06
MAIL=/var/mail/jschauma
HOME=/home/jschauma
PATH=/home/jschauma/bin:/bin:/sbin:/usr/bin:/usr/sbin:/usr/X11R7/bin:/usr/pkg/bin:/usr/pkg/sbin:/usr/games:/usr/local/bin:/usr/local/sbin
SSH_CONNECTION=10.0.2.2 61988 10.0.2.15 22
SSH_TTY=/dev/pts/1
TERM=xterm-256color
OLDPWD=/home/jschauma
USER=jschauma
SSH_AUTH_INFO_0=publickey ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQCV0TegkoPwNoeNf8
yWpc+N5Q2SS8Gz0pwQ88xsjElrPl4KsulC0n3EGu8Ry/SJdV8Ccfg0Wt21vMRajltC3sxxmCKjL0kAR
pPk++EVPu/WbkP5tsuKBPAPlaKUyfGiYkQtaPW04eP5Crr+c0qbyAqn+cHnAP/sMFuZvHPwhFK3ywaR/
K/o87hrlHg030rn7v4wDEn/a+TJMmpLZ12IQ06/PcfFuyKySuf+NXd/3m4dX9UPIzH5g4Y4HrEVjK4aE
Au9kAsd1SfkyUZUoRat0NAEanUm7hznzffJRrdMs1s2pixb0BdSU0QNy2E67Iwa70yf2fxMQDorQJFxJ
VuC5sN

CFLAGS=-Wall -Werror -Wextra
SSH_CLIENT=10.0.2.2 61988 22
EDITOR=vi
LOGNAME=jschauma
```

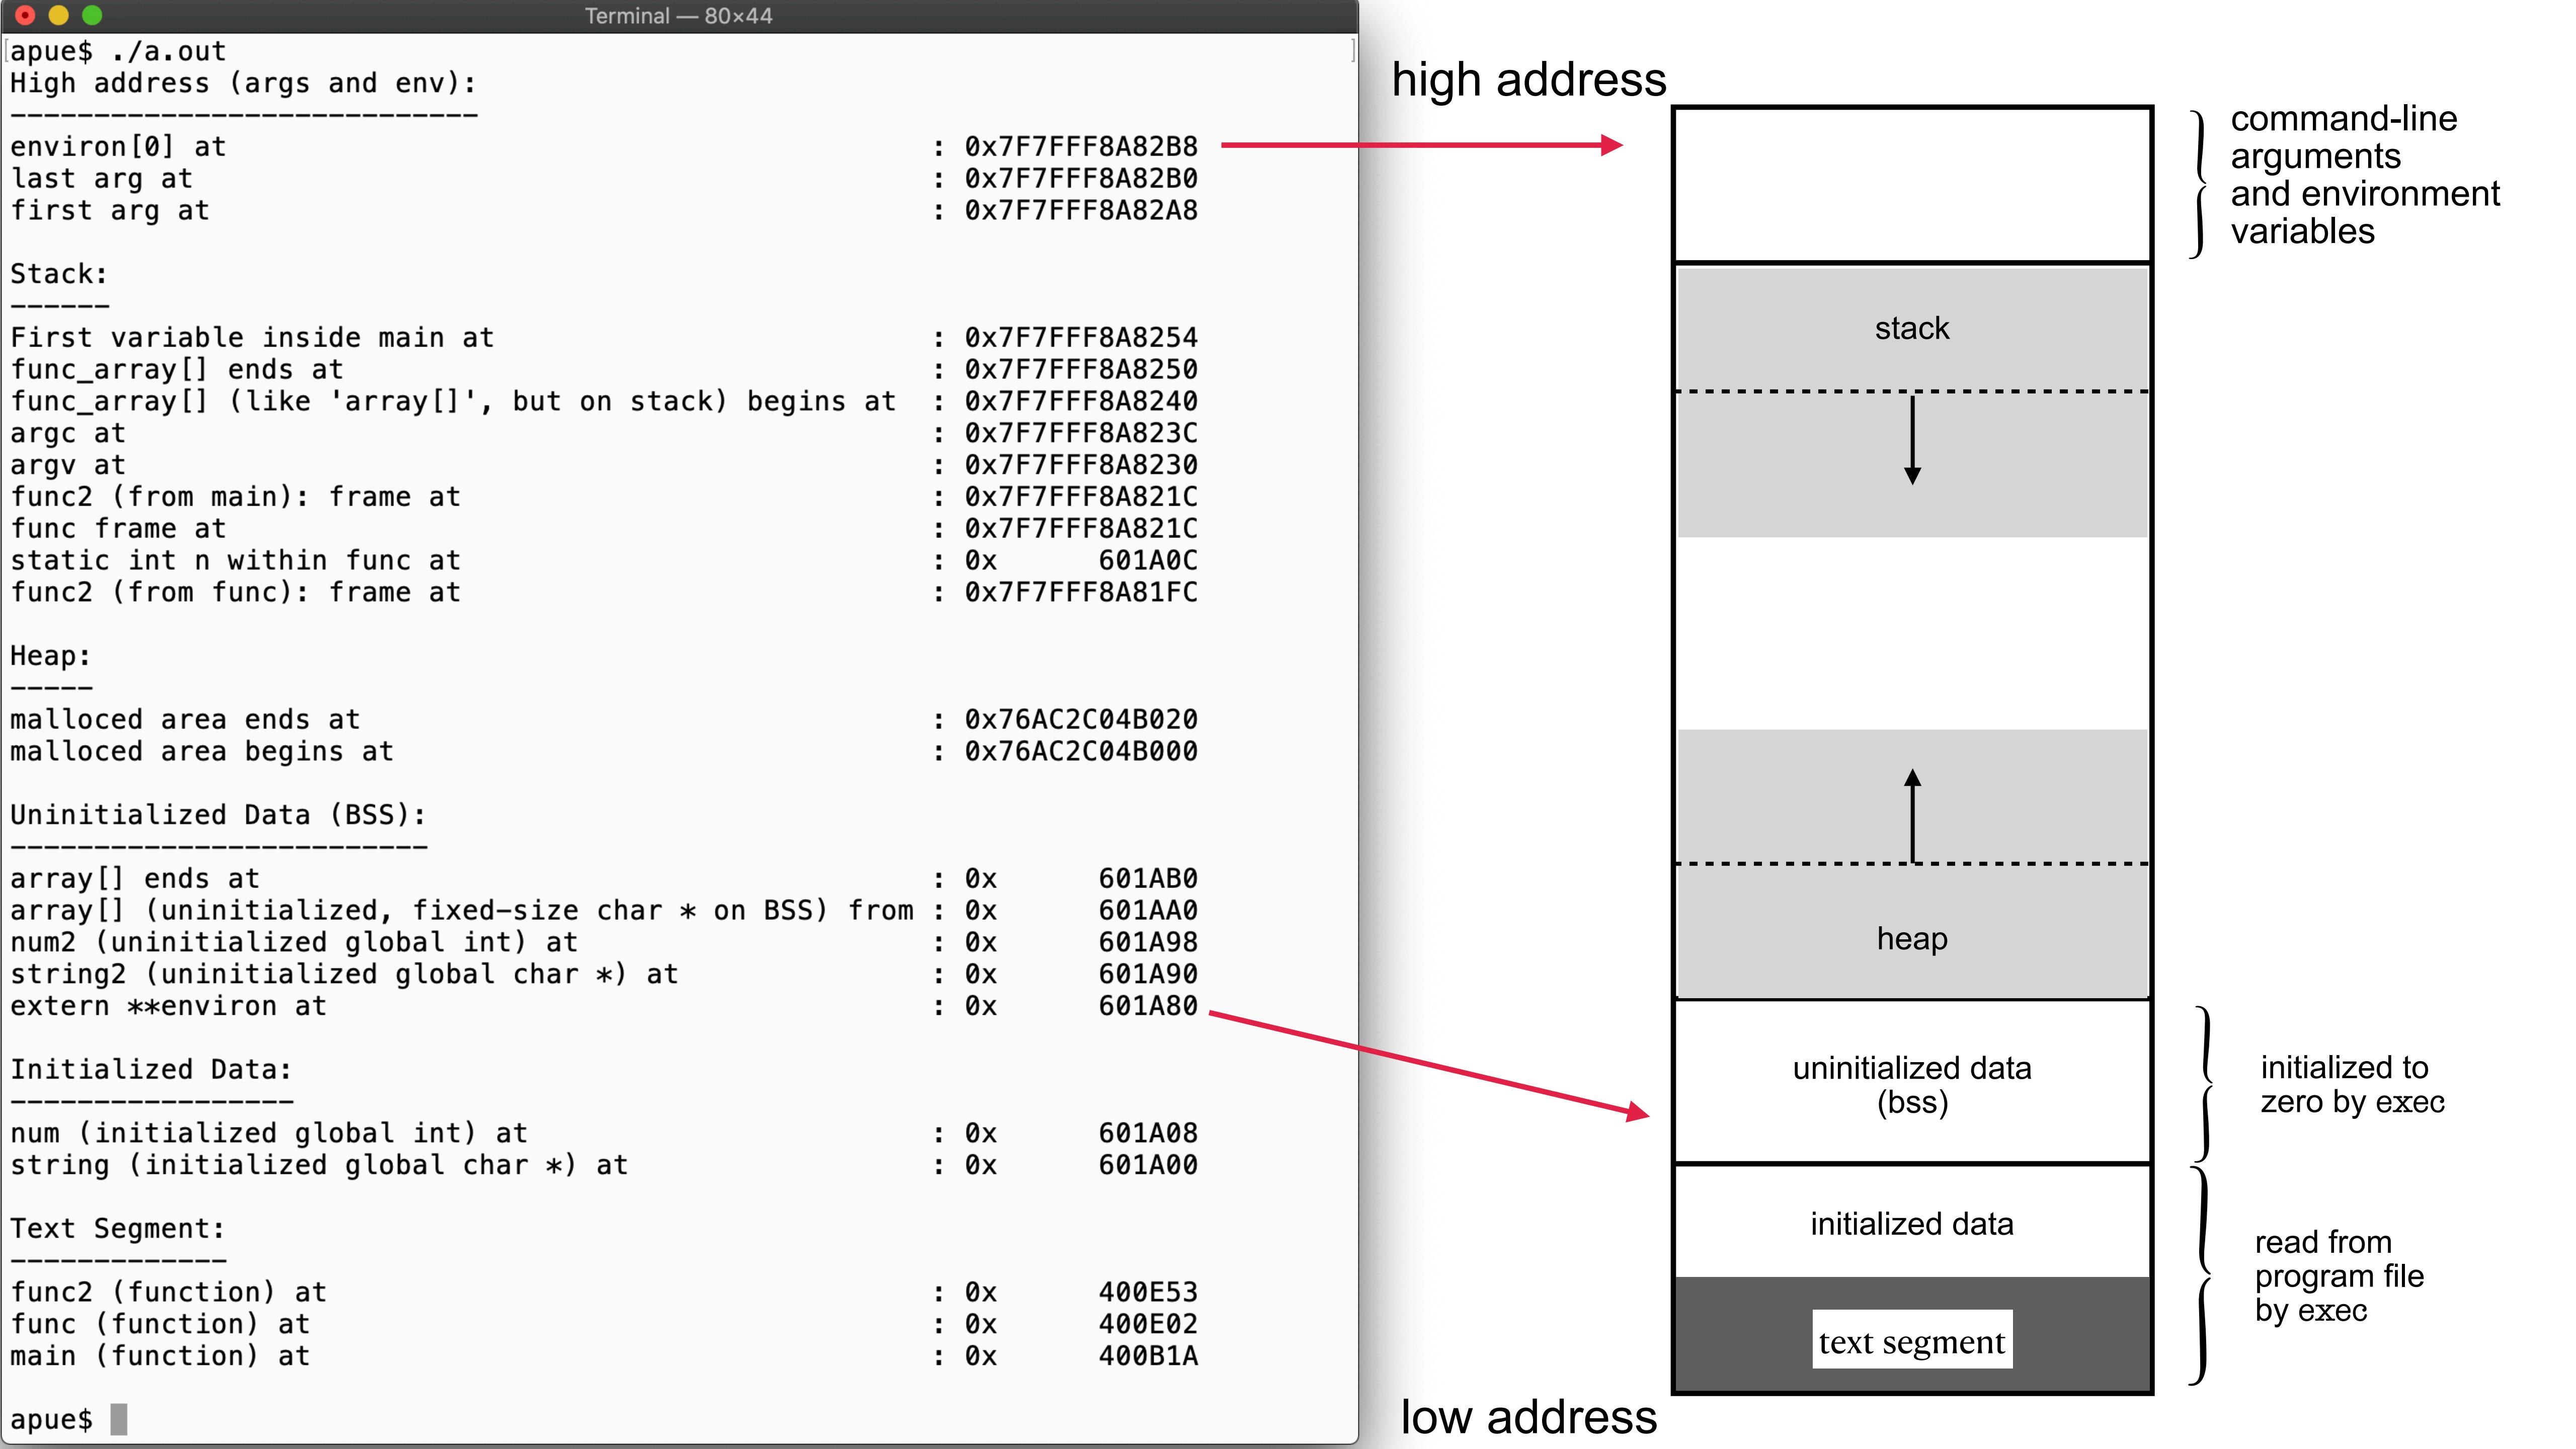
high address



low address

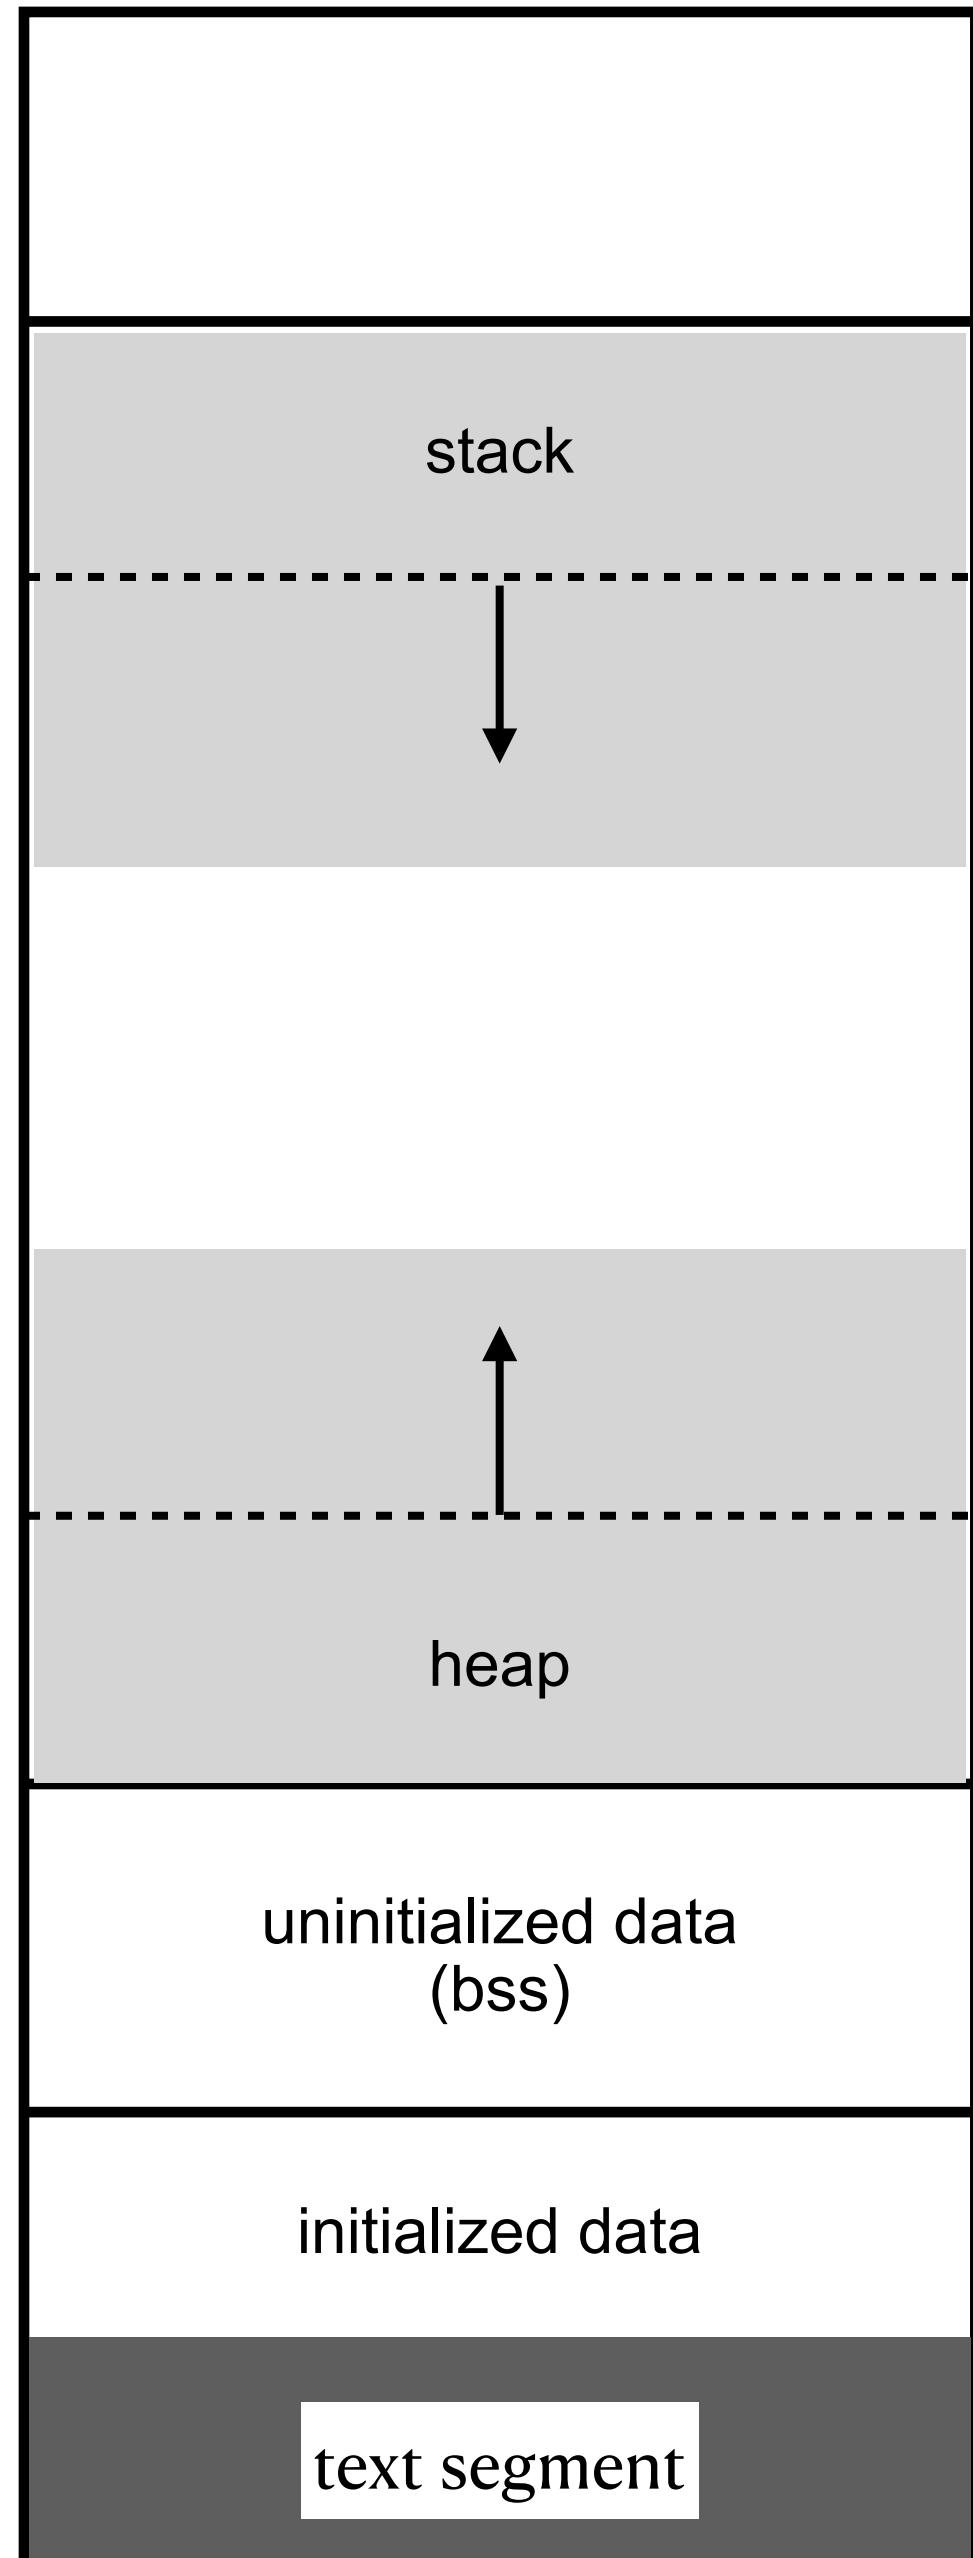
initialized to zero by exec

read from program file by exec



high address

low address



} command-line arguments and environment variables

} initialized to zero by exec

} read from program file by exec



```
relocate_self(ps_strings);
#endif

if (ps_strings == NULL)
    _FATAL("ps_strings missing\n");
__ps_strings = ps_strings;

environ = ps_strings->ps_envstr;

if (ps_strings->ps_argvstr[0] != NULL) {
    char *c;
    __progname = ps_strings->ps_argvstr[0];
    for (c = ps_strings->ps_argvstr[0]; *c; ++c) {
        if (*c == '/')
            __progname = c + 1;
    }
} else {
    __progname = empty_string;
}
--- 29 lines: if (cleanup != NULL)-----
exit(main(ps_strings->ps_nargvstr, ps_strings->ps_argvstr, environ));
}
/environ
```

Environment List

Environment variables are stored in a global, NULL terminated array of pointers:

```
extern char **environ;
```

A similar array may also have been passed into your main function:

```
int main(int argc, char **argv, char **envp);
```



```
--- memory-layout3.c      2020-10-01 02:44:07.242554206 +0000
+++ memory-layout4.c      2020-10-04 18:34:32.156169438 +0000
@@ -30,4 +30,4 @@
int
-main(int argc, char **argv) {
-    int var;
+main(int argc, char **argv, char **envp) {
+    int vars;
        char *ptr;
@@ -36,4 +36,13 @@
+
+    vars = 0;
+    char **tmp = envp;
+    while (*tmp++) {
+        vars++;
+
+        (void)printf("High address (args and env):\n");
+        (void)printf("-----\n");
+        (void)printf("envp[%d] at
x%12lX\n", vars, (unsigned long)&envp[vars]);
+        (void)printf("environ[%d] at
x%12lX\n", vars, (unsigned long)&environ[vars]);
--More--(byte 681)
```

```
[apue$ ./a.out
High address (args and env):
```

```
envp[17] at
environ[17] at
envp[0] at
environ[0] at
last arg at
first arg at
```

Stack:

```
First variable inside main at
func_array[] ends at
func_array[] (like 'array[]', but on stack) begins at
argc at
argv at
envp at
func2 (from main): frame at
func frame at
static int n within func at
func (called 1 times): frame at
func2 (from func): frame at
```

Heap:

```
malloced area ends at
malloced area begins at
```

Uninitialized Data (BSS):

```
array[] ends at : 0x 601CF0
array[] (uninitialized, fixed-size char * on BSS) from : 0x 601CE0
num2 (uninitialized global int) at : 0x 601CD8
string2 (uninitialized global char *) at : 0x 601CD0
extern **environ at : 0x 601CC0
```

Initialized Data:

```
num (initialized global int) at : 0x 601C60
string (initialized global char *) at : 0x 601C58
```

Text Segment:

```
func2 (function) at : 0x 400F3F
func (function) at : 0x 400EB8
main (function) at : 0x 400B1A
```

high address

low address

```
: 0x7FFF66D860
: 0x7FFF66D860
: 0x7FFF66D7D8
: 0x7FFF66D7D8
: 0x7FFF66D7D0
: 0x7FFF66D7C8
```

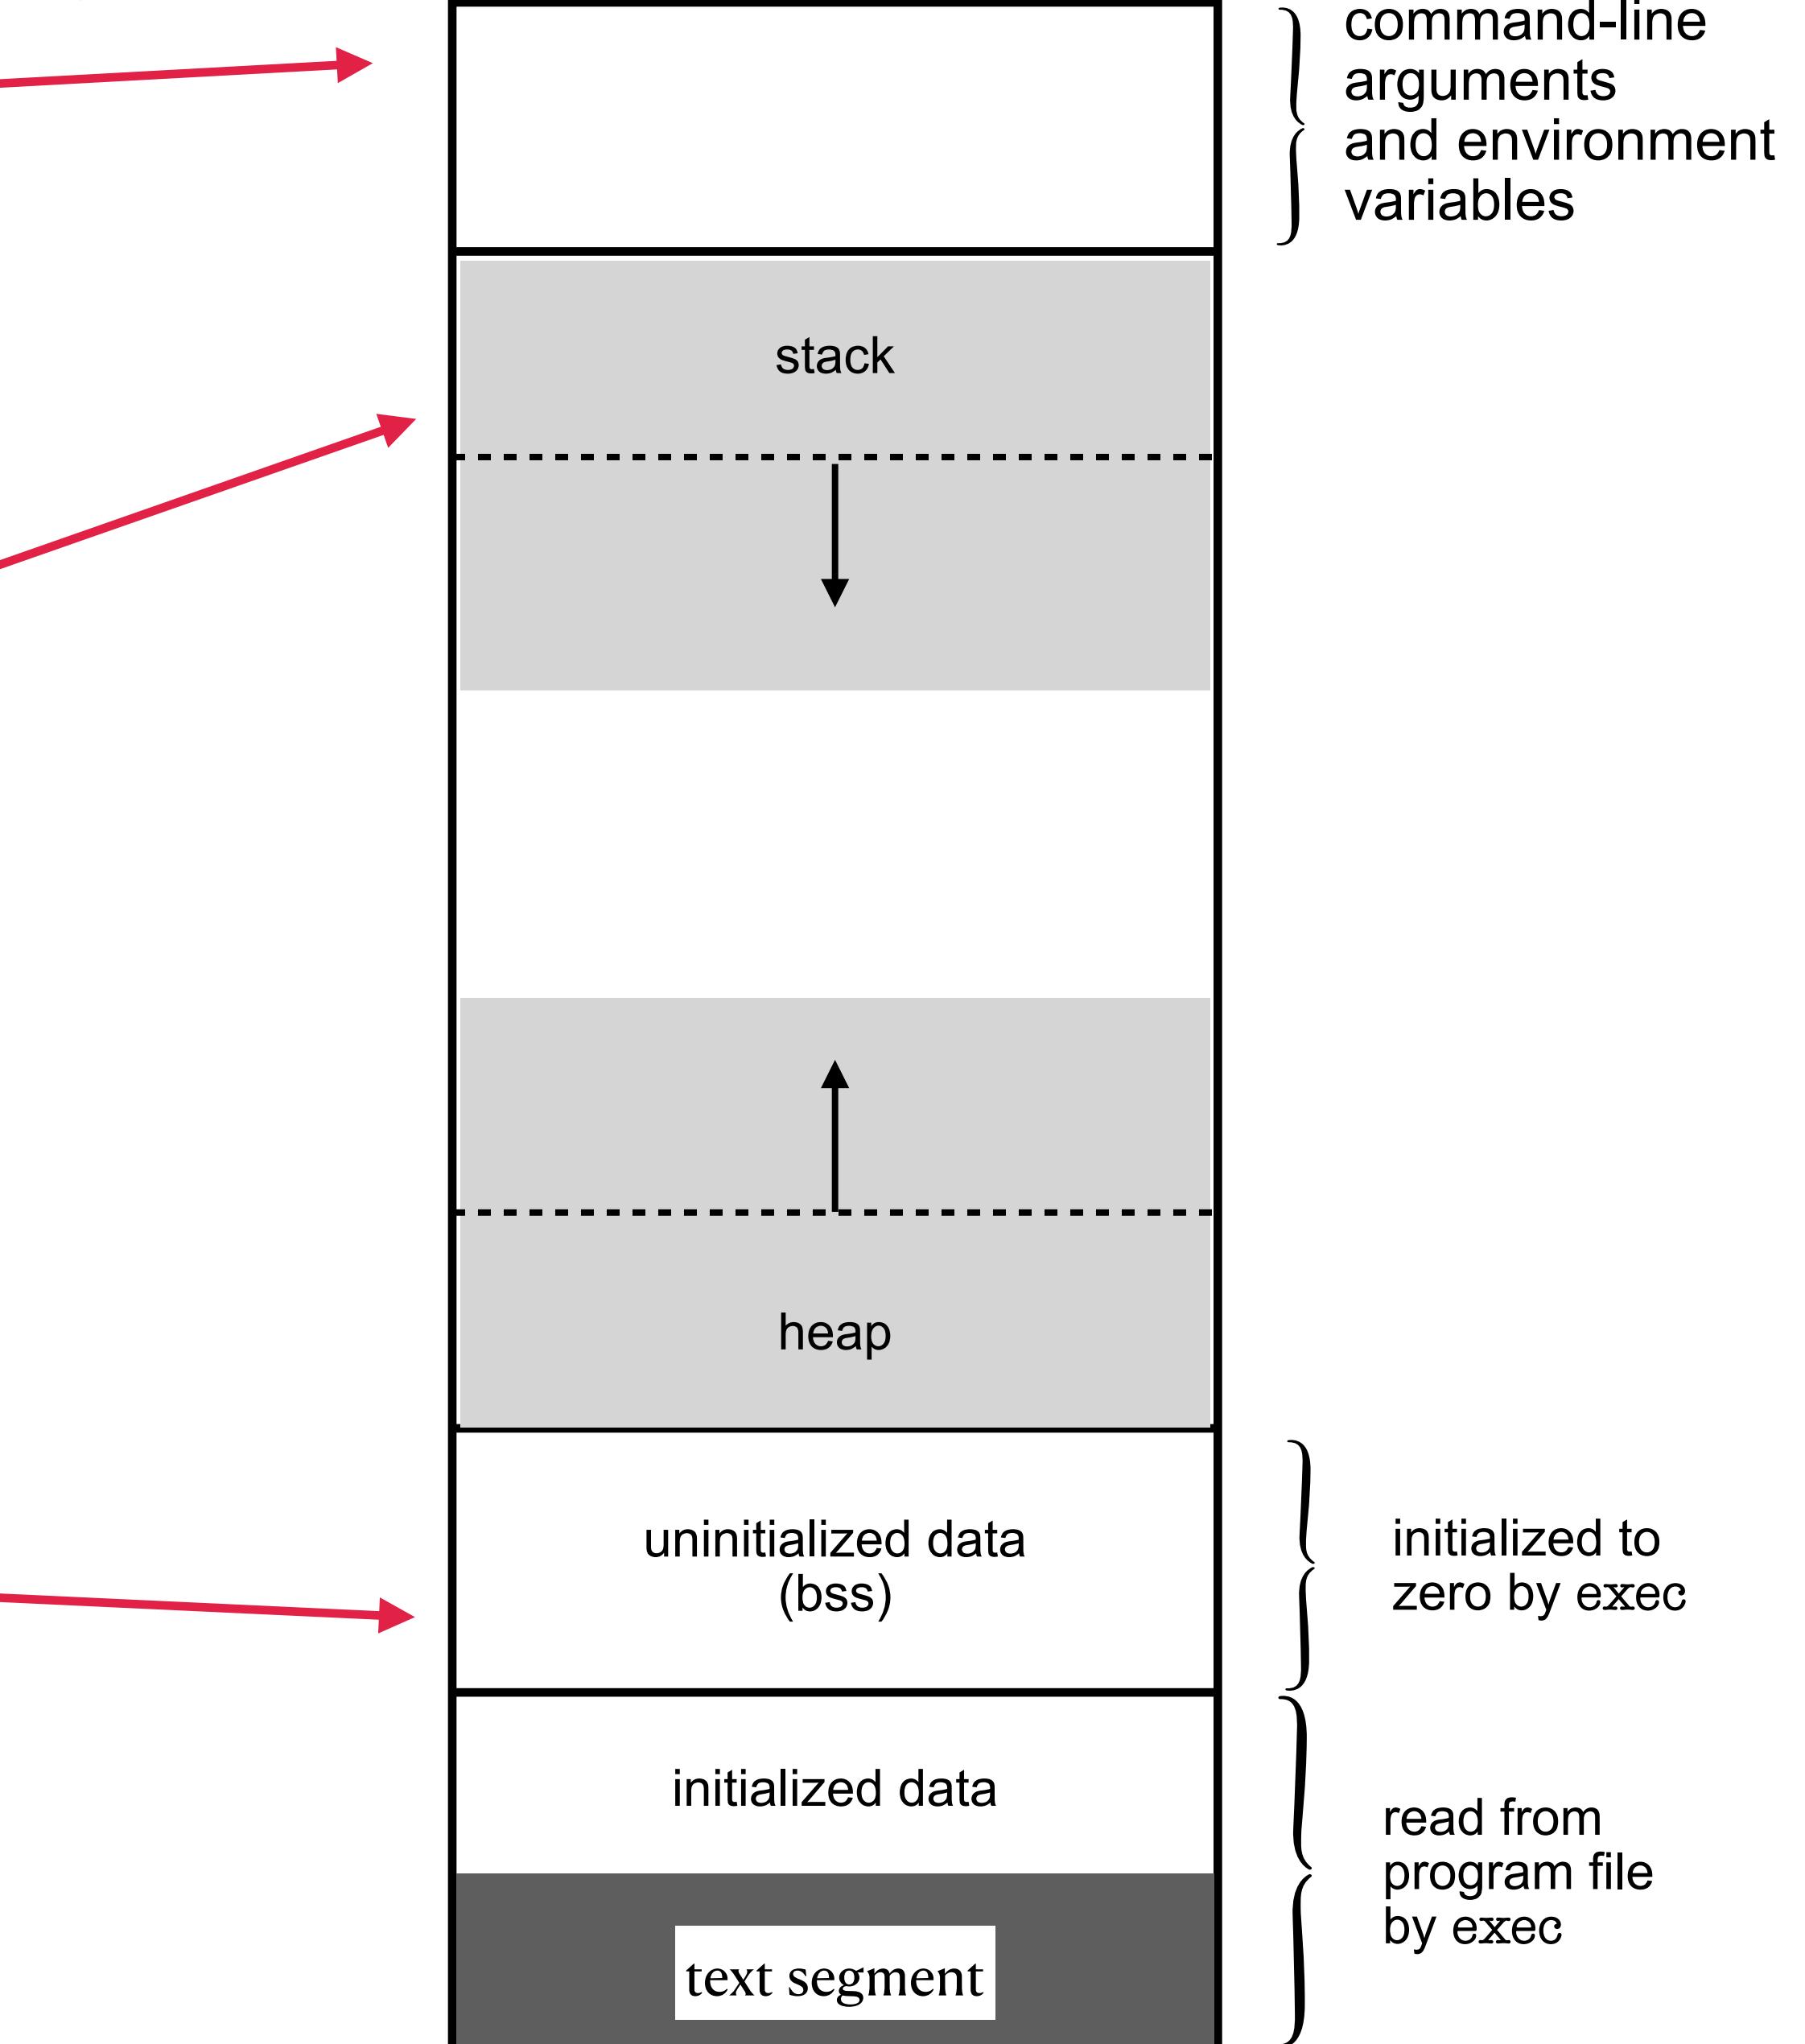
```
: 0x7FFF66D76C
: 0x7FFF66D760
: 0x7FFF66D750
: 0x7FFF66D74C
: 0x7FFF66D740
: 0x7FFF66D738
: 0x7FFF66D71C
: 0x7FFF66D714
: 0x 601C64
: 0x7FFF66D714
: 0x7FFF66D6EC
```

```
: 0x759D760BF020
: 0x759D760BF000
```

```
: 0x 601CF0
: 0x 601CE0
: 0x 601CD8
: 0x 601CD0
: 0x 601CC0
```

```
: 0x 601C60
: 0x 601C58
```

```
: 0x 400F3F
: 0x 400EB8
: 0x 400B1A
```



} command-line
arguments
and environment
variables

} initialized to
zero by exec

} read from
program file
by exec

getenv(3) etc.

```
#include <stdlib.h>  
char *getenv(const char *name);
```

Returns: value if found, NULL otherwise

```
int putenv(char *string);
```

```
int setenv(const char *name, const char *value, int overwrite);
```

```
int unsetenv(const char *name);
```

Returns: 0 on success; -1 on error

Note the difference between unsetenv("VAR") and setenv("VAR", "", 1)!

```

Terminal — 80x44
[apue$ cc -Wall -Werror -Wextra memory-layout5.c
[apue$ ./a.out
At program start:
0x7F7FFFBCBBA8  envp [18]
0x7F7FFFBCBBA0  envp [17] (0x7F7FFFBCC40B 'SHELL=/bin/sh')
0x7F7FFFBCBB18  envp [0]  (0x7F7FFFBCC0B0 'FOO=bar')
...
0x7F7FFFBCBBA8  environ[18]
0x7F7FFFBCBBA0  environ[17] (0x7F7FFFBCC40B 'SHELL=/bin/sh')
0x7F7FFFBCBB18  environ[0]  (0x7F7FFFBCC0B0 'FOO=bar')
...
Stack:
-----
0x7F7FFFBCBAAC int argc
0x7F7FFFBCBAA0 char **argv
0x7F7FFFBCBA98 char **envp

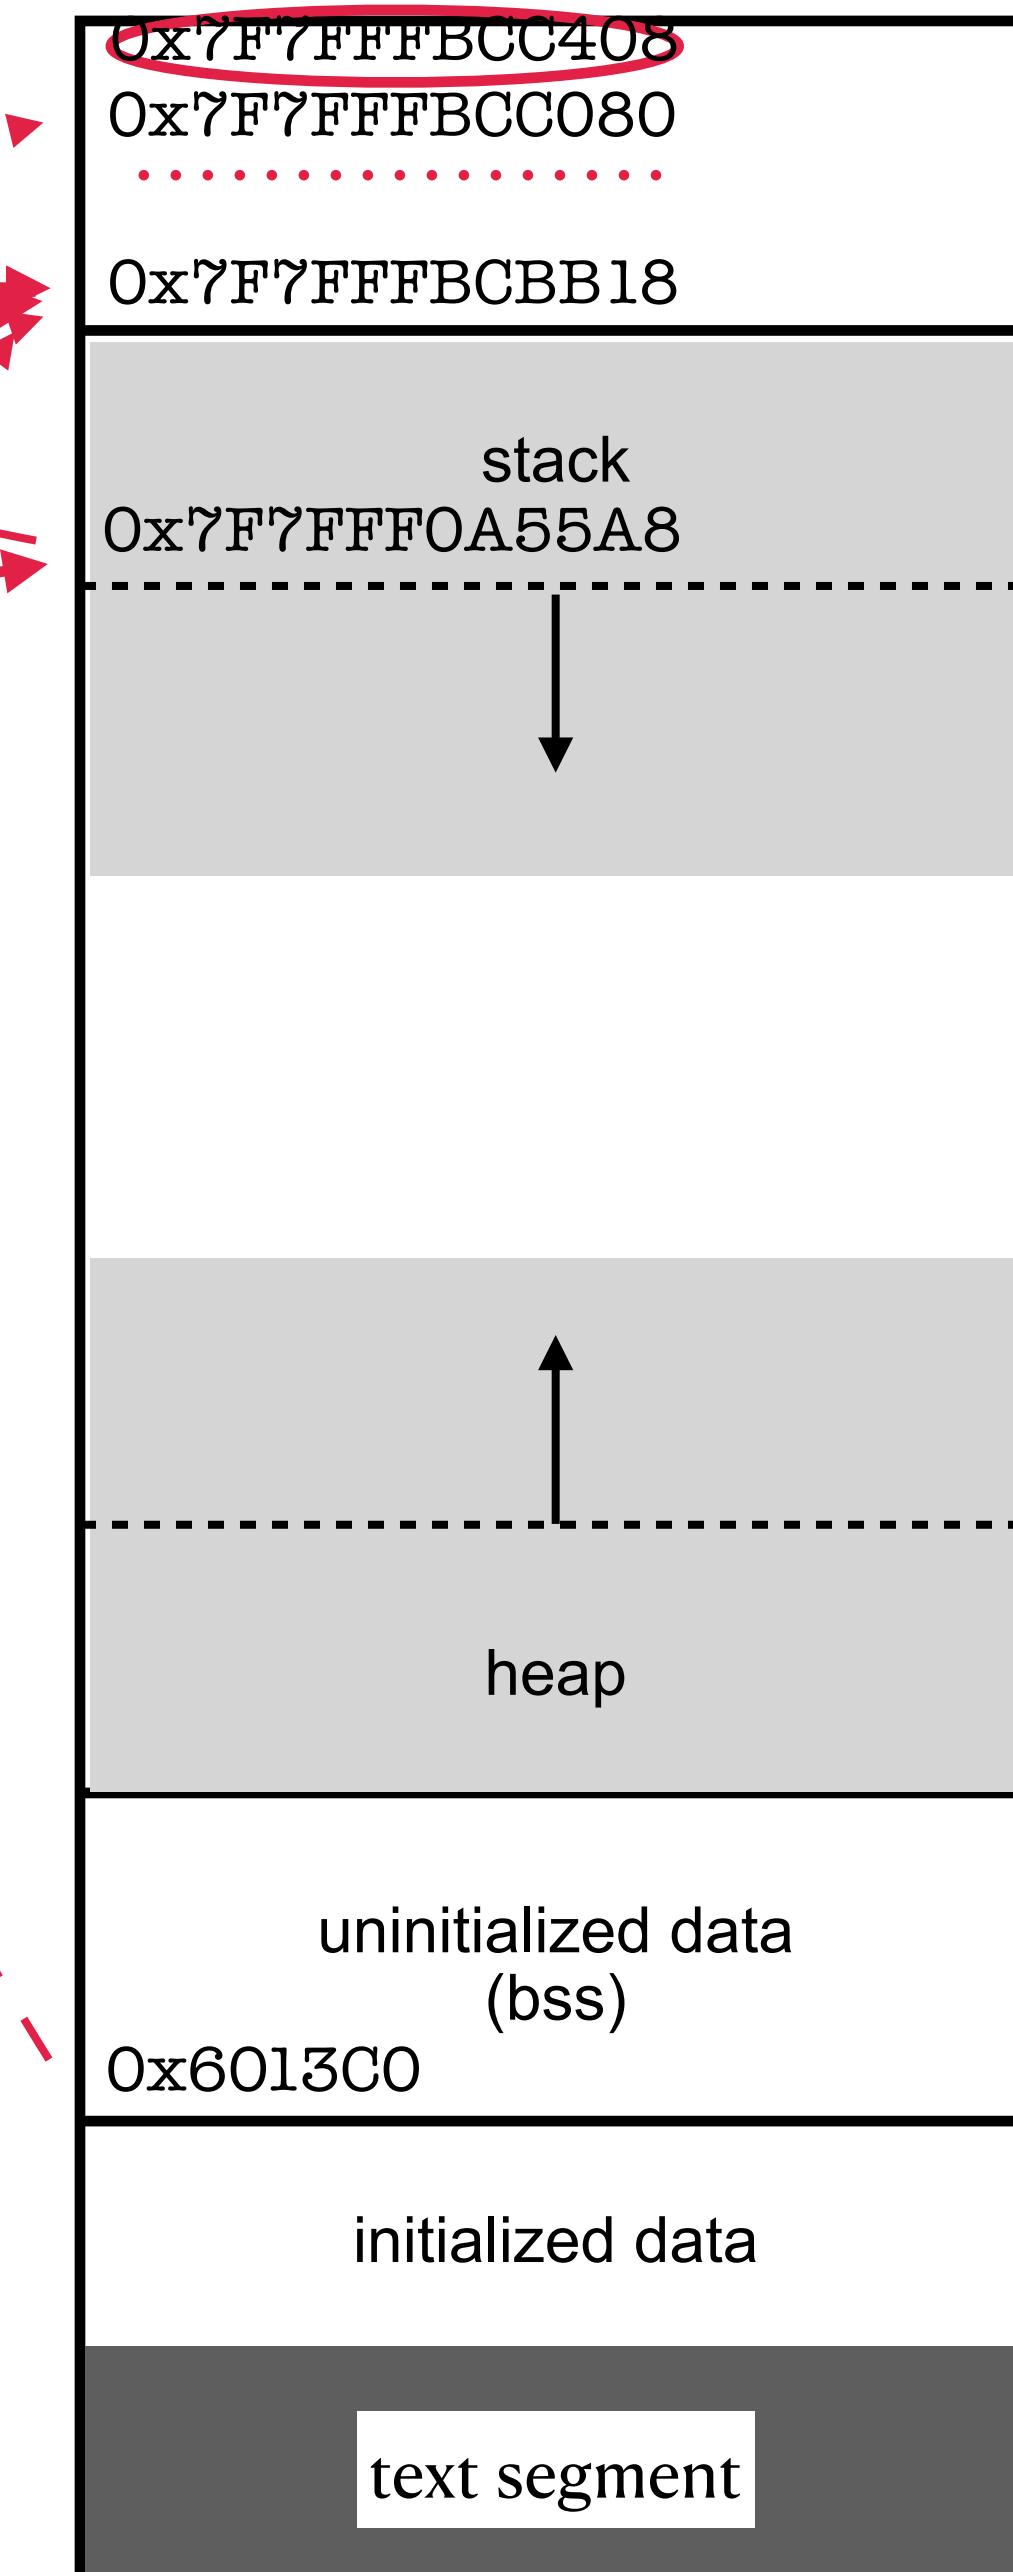
Heap:
-----
0x77B13E0BD020 malloced area ends
0x77B13E0BD000 malloced area begins

Uninitialized Data (BSS):
-----
0x 6013C0 extern char **environ
apue$ 

```

high address

low address



} command-line
arguments
and environment
variables

} initialized to
zero by exec

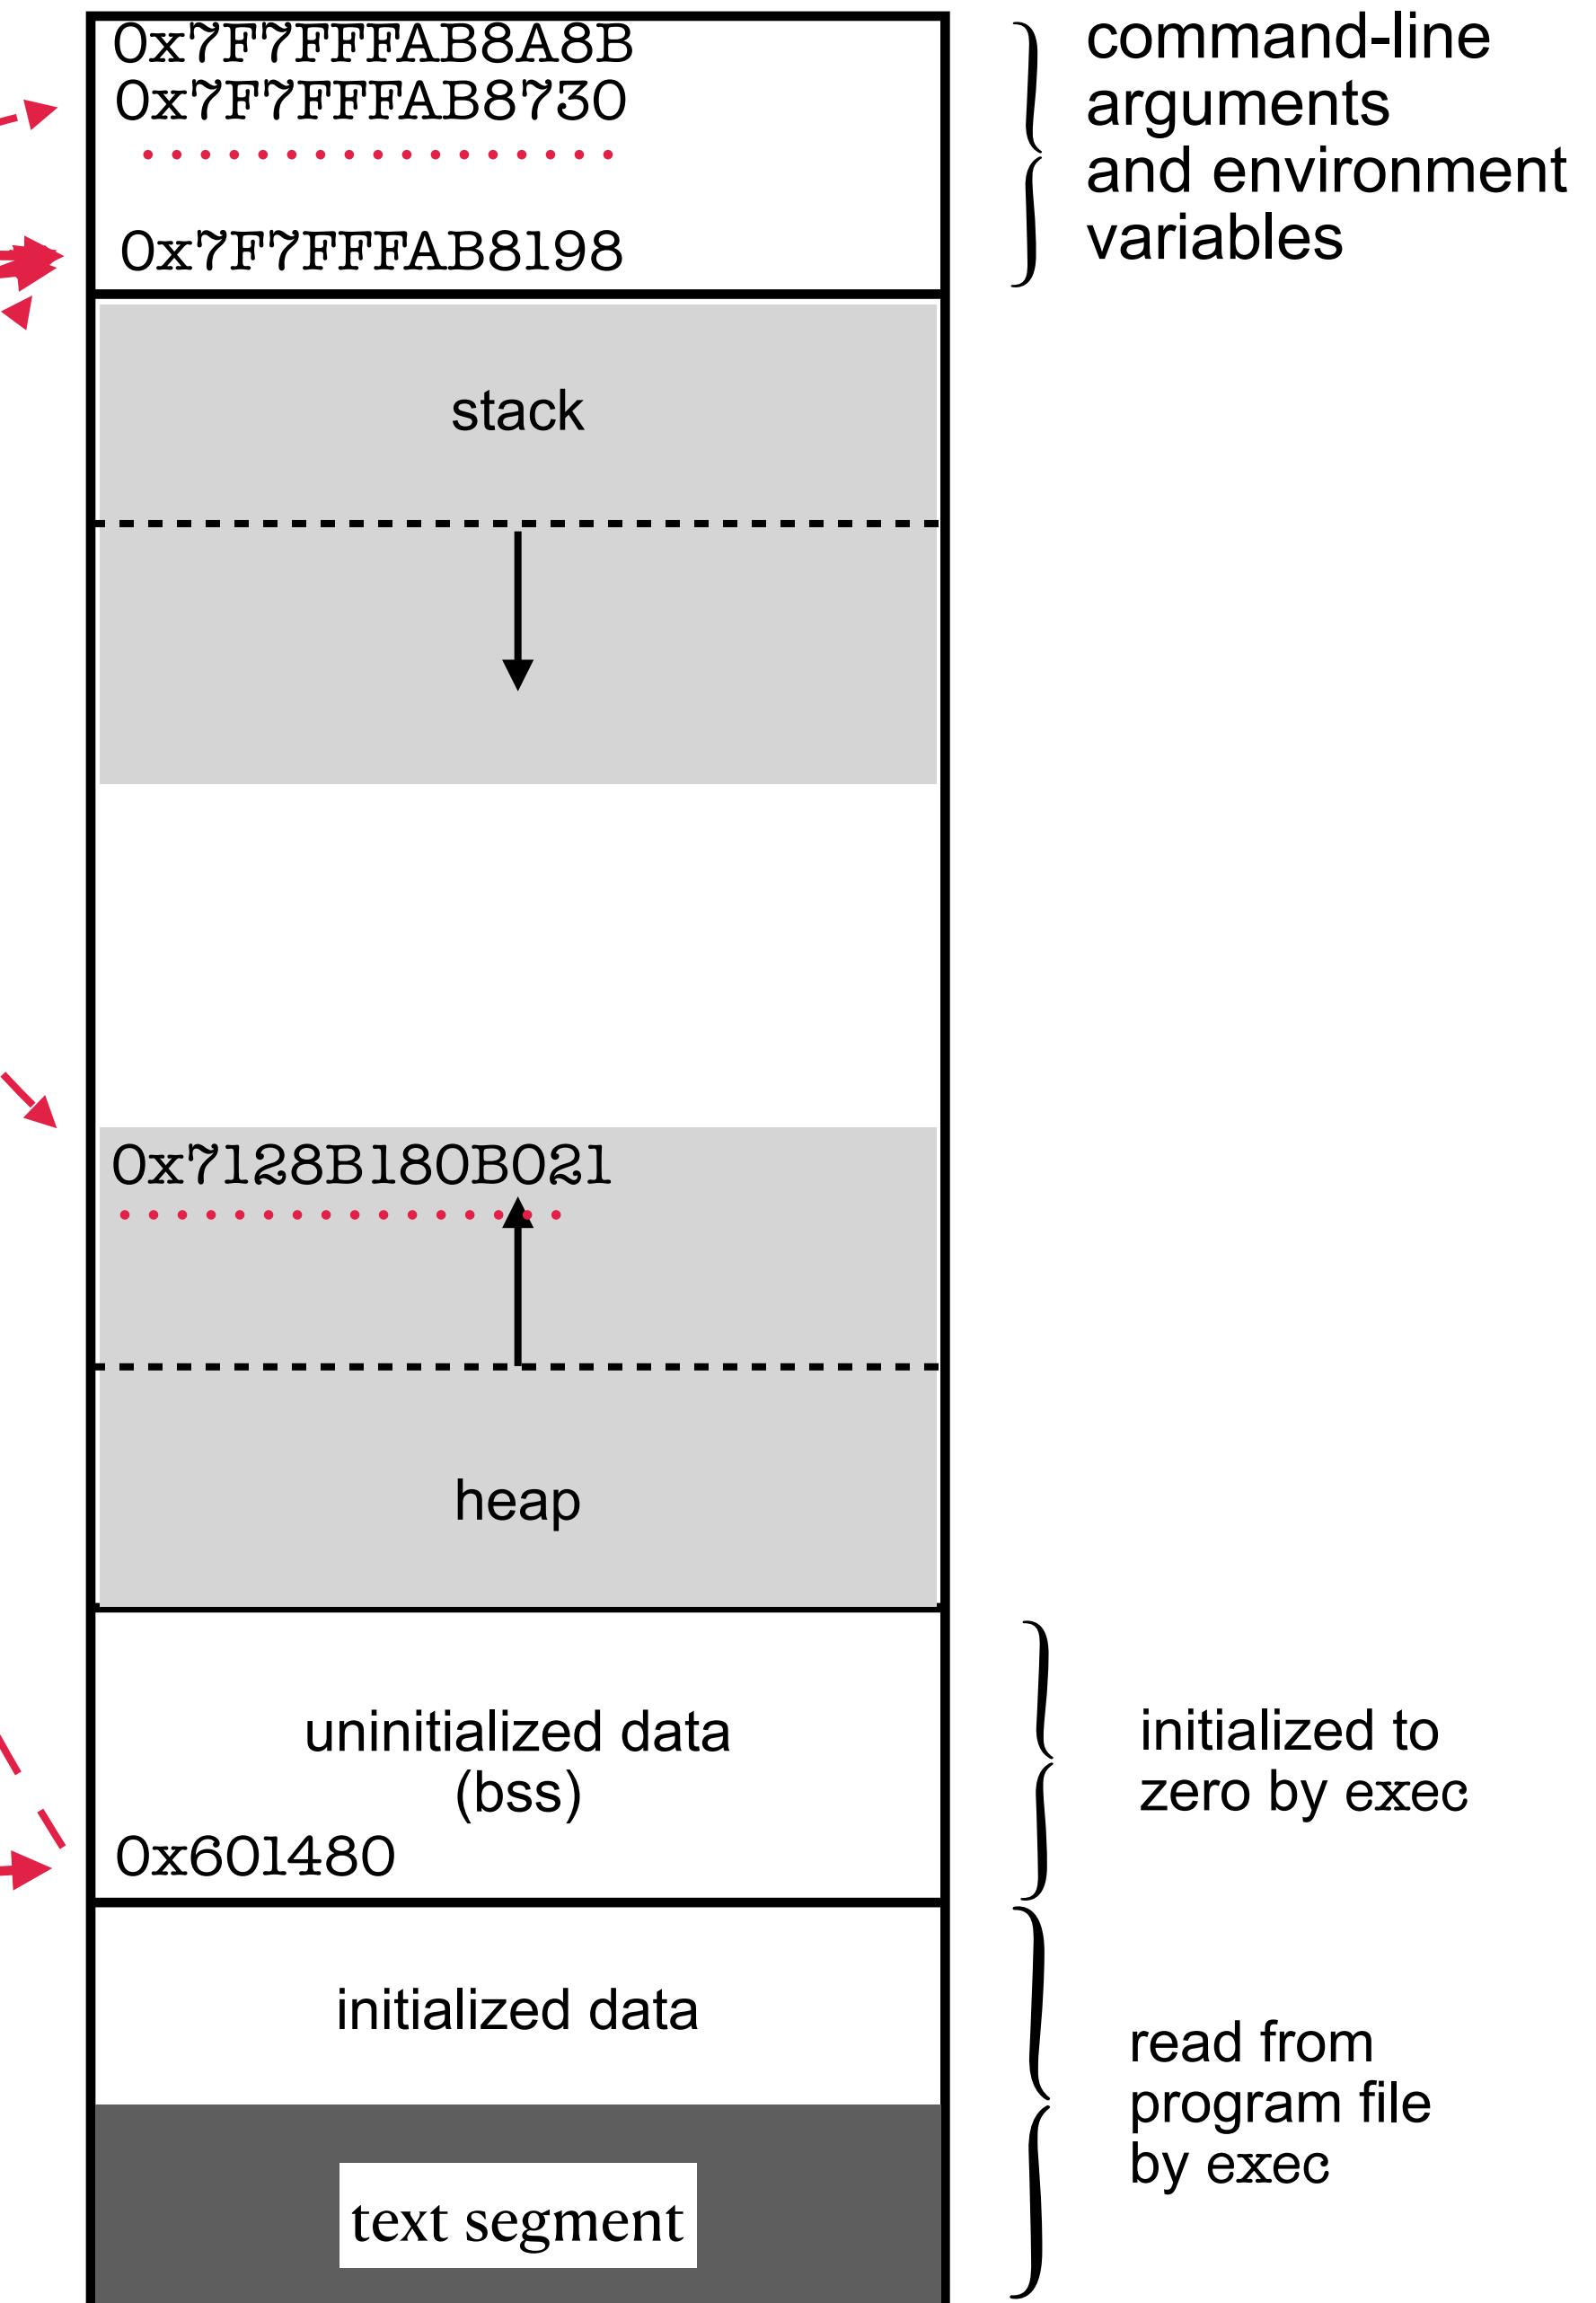
} read from
program file
by exec

Terminal — 80x44

```
[apue$ cc -Wall -Werror -Wextra memory-layout6.c
[apue$ ./a.out
At program start:
0x7F7FFFAB8228  envp  [18]
0x7F7FFFAB8220  envp  [17]  (0x7F7FFFAB8A8B 'SHELL=/bin/sh')
0x7F7FFFAB8198  envp  [0]  (0x7F7FFFAB8730 'FOO=bar')
...
0x7F7FFFAB8228  environ[18]
0x7F7FFFAB8220  environ[17]  (0x7F7FFFAB8A8B 'SHELL=/bin/sh')
0x7F7FFFAB8198  environ[0]  (0x7F7FFFAB8730 'FOO=bar')
...
After setenv(3):
0x7F7FFFAB8228  envp  [18]
0x7F7FFFAB8220  envp  [17]  (0x7F7FFFAB8A8B 'SHELL=/bin/sh')
0x7F7FFFAB8198  envp  [0]  (0x7128B180B021 'FOO=a longer value')
...
0x7F7FFFAB8228  environ[18]
0x7F7FFFAB8220  environ[17]  (0x7F7FFFAB8A8B 'SHELL=/bin/sh')
0x7F7FFFAB8198  environ[0]  (0x7128B180B021 'FOO=a longer value')
...
Stack:
-----
0x7F7FFFAB812C  int argc
0x7F7FFFAB8120  char **argv
0x7F7FFFAB8118  char **envp
...
Heap:
-----
0x7128B180A020  malloced area ends
0x7128B180A000  malloced area begins
...
Uninitialized Data (BSS):
-----
0x      6014C0  extern char **environ
...
apue$ ]
```

high address

low address



malloc(3) and friends

```
#include <stdlib.h>

void *malloc(size_t size);
void *calloc(size_t number, size_t size);
void *realloc(void *ptr, size_t size);

void free(void *ptr);
```

Returns: pointer on success, NULL otherwise

malloc(3) allocates size bytes of uninitialized memory.

calloc(3) allocates number * size bytes of memory, initialized to zero bytes.

realloc(3) lets you grow or shrink a previously allocated area, possibly returning a new pointer.

See jemalloc(3) for a lot more details.

```
Terminal — 65x44
#define MALLOC_SIZE 1024 * 1024

void
printMalloc(void *p, int s) {
    (void)printf("malloc/realloc %d\n", s);
    (void)printf("begins at 0x%12lX\n", (unsigned long)p);
    (void)printf("ends   at 0x%12lX\n", (unsigned long)p+s);
    (void)printf("\n");
}

int
main() {
    void *ptr;

    if ((ptr = malloc(BUFSIZ)) == NULL) {
        err(EXIT_FAILURE, "malloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ);

    /* shrink */
    if ((ptr = realloc(ptr, BUFSIZ / 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ / 2);

    /* grow */
    if ((ptr = realloc(ptr, MALLOC_SIZE)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, MALLOC_SIZE);

    /* shrink, but larger than initial allocation */
    if ((ptr = realloc(ptr, BUFSIZ * 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ * 4);

    return EXIT_SUCCESS;
}
apue$
```

```
Terminal — 65x44
[apue$ cc -Wall -Werror -Wextra malloc.c
[apue$ ./a.out
malloc/realloc 1024
begins at 0x71D072422000
ends   at 0x71D072422400

malloc/realloc 512
begins at 0x71D072415000
ends   at 0x71D072415200

malloc/realloc 1048576
begins at 0x71D0716FF2C0
ends   at 0x71D0717FF2C0

malloc/realloc 4096
begins at 0x71D072413000
ends   at 0x71D072414000

apue$
```

The diagram illustrates the memory layout of the process. It shows four distinct segments: 'stack' at the top, followed by a dashed line, then 'uninitialized data (bss)', 'initialized data', and finally 'text segment' at the bottom. A red arrow points from the terminal output 'begins at 0x71D072422400' to the 'uninitialized data (bss)' segment. Another red arrow points from the terminal output 'ends at 0x71D072422000' to the boundary between the 'uninitialized data (bss)' and 'initialized data' segments.

```
Terminal — 65x44
#define MALLOC_SIZE 1024 * 1024

void
printMalloc(void *p, int s) {
    (void)printf("malloc/realloc %d\n", s);
    (void)printf("begins at 0x%12lX\n", (unsigned long)p);
    (void)printf("ends   at 0x%12lX\n", (unsigned long)p+s);
    (void)printf("\n");
}

int
main() {
    void *ptr;

    if ((ptr = malloc(BUFSIZ)) == NULL) {
        err(EXIT_FAILURE, "malloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ);

    /* shrink */
    if ((ptr = realloc(ptr, BUFSIZ / 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ / 2);

    /* grow */
    if ((ptr = realloc(ptr, MALLOC_SIZE)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, MALLOC_SIZE);

    /* shrink, but larger than initial allocation */
    if ((ptr = realloc(ptr, BUFSIZ * 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ * 4);

    return EXIT_SUCCESS;
}
apue$
```

```
Terminal — 65x44
[apue$ cc -Wall -Werror -Wextra malloc.c
[apue$ ./a.out
malloc/realloc 1024
begins at 0x71D072422000
ends   at 0x71D072422400

malloc/realloc 512
begins at 0x71D072415000
ends   at 0x71D072415200

malloc/realloc 1048576
begins at 0x71D0716FF2C0
ends   at 0x71D0717FF2C0

malloc/realloc 4096
begins at 0x71D072413000
ends   at 0x71D072414000

apue$
```

The diagram illustrates the memory layout of the process. It shows a vertical stack of memory regions. At the top is the 'stack' region, indicated by a downward arrow. Below it is the 'uninitialized data (bss)' region, indicated by an upward arrow. The 'initialized data' region is shown as a grey bar. At the bottom is the 'text segment' region, shown as a dark grey bar. A red arrow points from the highlighted memory addresses in the terminal output to the corresponding regions in the memory diagram.

stack

uninitialized data (bss)

initialized data

text segment

```
Terminal — 65x44
#define MALLOC_SIZE 1024 * 1024

void
printMalloc(void *p, int s) {
    (void)printf("malloc/realloc %d\n", s);
    (void)printf("begins at 0x%12lX\n", (unsigned long)p);
    (void)printf("ends   at 0x%12lX\n", (unsigned long)p+s);
    (void)printf("\n");
}

int
main() {
    void *ptr;

    if ((ptr = malloc(BUFSIZ)) == NULL) {
        err(EXIT_FAILURE, "malloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ);

    /* shrink */
    if ((ptr = realloc(ptr, BUFSIZ / 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ / 2);

    /* grow */
    if ((ptr = realloc(ptr, MALLOC_SIZE)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, MALLOC_SIZE);

    /* shrink, but larger than initial allocation */
    if ((ptr = realloc(ptr, BUFSIZ * 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ * 4);

    return EXIT_SUCCESS;
}
apue$
```

```
Terminal — 65x44
[apue$ cc -Wall -Werror -Wextra malloc.c
[apue$ ./a.out
malloc/realloc 1024
begins at 0x71D072422000
ends   at 0x71D072422400

malloc/realloc 512
begins at 0x71D072415000
ends   at 0x71D072415200

malloc/realloc 1048576
begins at 0x71D0716FF2C0
ends   at 0x71D0717FF2C0
apue$
```

The diagram illustrates the memory layout of the process. It shows a vertical stack of memory segments. At the top is the 'stack' segment, indicated by a downward-pointing arrow. Below it is the 'initialized data' segment, indicated by an upward-pointing arrow. The 'uninitialized data (bss)' segment is shown as a large block below the initialized data. The 'text segment' is at the bottom. Two specific memory addresses are highlighted with red boxes: 0x71D072415200 and 0x71D0716FF2C0. A red arrow points from the text 'begins at 0x71D0716FF2C0' in the terminal output to the address 0x71D0716FF2C0 in the memory diagram. Another red arrow points from the text 'ends at 0x71D0717FF2C0' in the terminal output to the address 0x71D0717FF2C0 in the memory diagram.

```
Terminal — 65x44
#define MALLOC_SIZE 1024 * 1024

void
printMalloc(void *p, int s) {
    (void)printf("malloc/realloc %d\n", s);
    (void)printf("begins at 0x%12lX\n", (unsigned long)p);
    (void)printf("ends   at 0x%12lX\n", (unsigned long)p+s);
    (void)printf("\n");
}

int
main() {
    void *ptr;

    if ((ptr = malloc(BUFSIZ)) == NULL) {
        err(EXIT_FAILURE, "malloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ);

    /* shrink */
    if ((ptr = realloc(ptr, BUFSIZ / 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ / 2);

    /* grow */
    if ((ptr = realloc(ptr, MALLOC_SIZE)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, MALLOC_SIZE);

    /* shrink, but larger than initial allocation */
    if ((ptr = realloc(ptr, BUFSIZ * 2)) == NULL) {
        err(EXIT_FAILURE, "realloc");
        /* NOTREACHED */
    }
    printMalloc(ptr, BUFSIZ * 4);

    return EXIT_SUCCESS;
}
apue$
```

```
Terminal — 65x44
[apue$ cc -Wall -Werror -Wextra malloc.c
[apue$ ./a.out
malloc/realloc 1024
begins at 0x71D072422000
ends   at 0x71D072422400

malloc/realloc 512
begins at 0x71D072415000
ends   at 0x71D072415200

malloc/realloc 1048576
begins at 0x71D0716FF2C0
ends   at 0x71D0717FF2C0

malloc/realloc 4096
begins at 0x71D072413000
ends   at 0x71D072414000
apue$
```

The diagram illustrates the memory layout of the process. The stack grows downwards from the top. The heap consists of several allocations:

- Allocation 1: begins at 0x71D072413000, ends at 0x71D072414000 (4096 bytes)
- Allocation 2: begins at 0x71D0717FF2C0, ends at 0x71D0717FF2C0 (1048576 bytes)
- Allocation 3: begins at 0x71D0716FF2C0, ends at 0x71D0716FF2C0 (512 bytes)
- Allocation 4: begins at 0x71D072422000, ends at 0x71D072422400 (1024 bytes)

Red boxes highlight the first two allocations (4096 and 1048576 bytes) in the dump, and a red arrow points from the 4096 byte allocation in the dump to its corresponding entry in the memory layout.

Terminal — 80x44

```
[apue$ cc -Wall -Werror -Wextra memory-layout7.c
[apue$ ./a.out
At program start:
0x7FFF666378  envp [18]          (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF666370  envp [17]          (0x7FFF666880 'FOO=bar')
0x7FFF6662E8  envp [0]           (0x7FFF666880 'FOO=bar')

0x7FFF666378  environ[18]        (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF666370  environ[17]        (0x7FFF666880 'FOO=bar')
0x7FFF6662E8  environ[0]         (0x7FFF666880 'FOO=bar')

After setenv(3):
0x7FFF666378  envp [18]          (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF666370  envp [17]          (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF6662E8  envp [0]           (0x78FDA87F4021 'FOO=a longer value')

0x7FFF666378  environ[18]        (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF666370  environ[17]        (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF6662E8  environ[0]         (0x78FDA87F4021 'FOO=a longer value')

After putenv(3):
0x7FFF666378  envp [18]          (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF666370  envp [17]          (0x7FFF666BDB 'SHELL=/bin/sh')
0x7FFF6662E8  envp [0]           (0x78FDA87F4021 'FOO=a longer value')

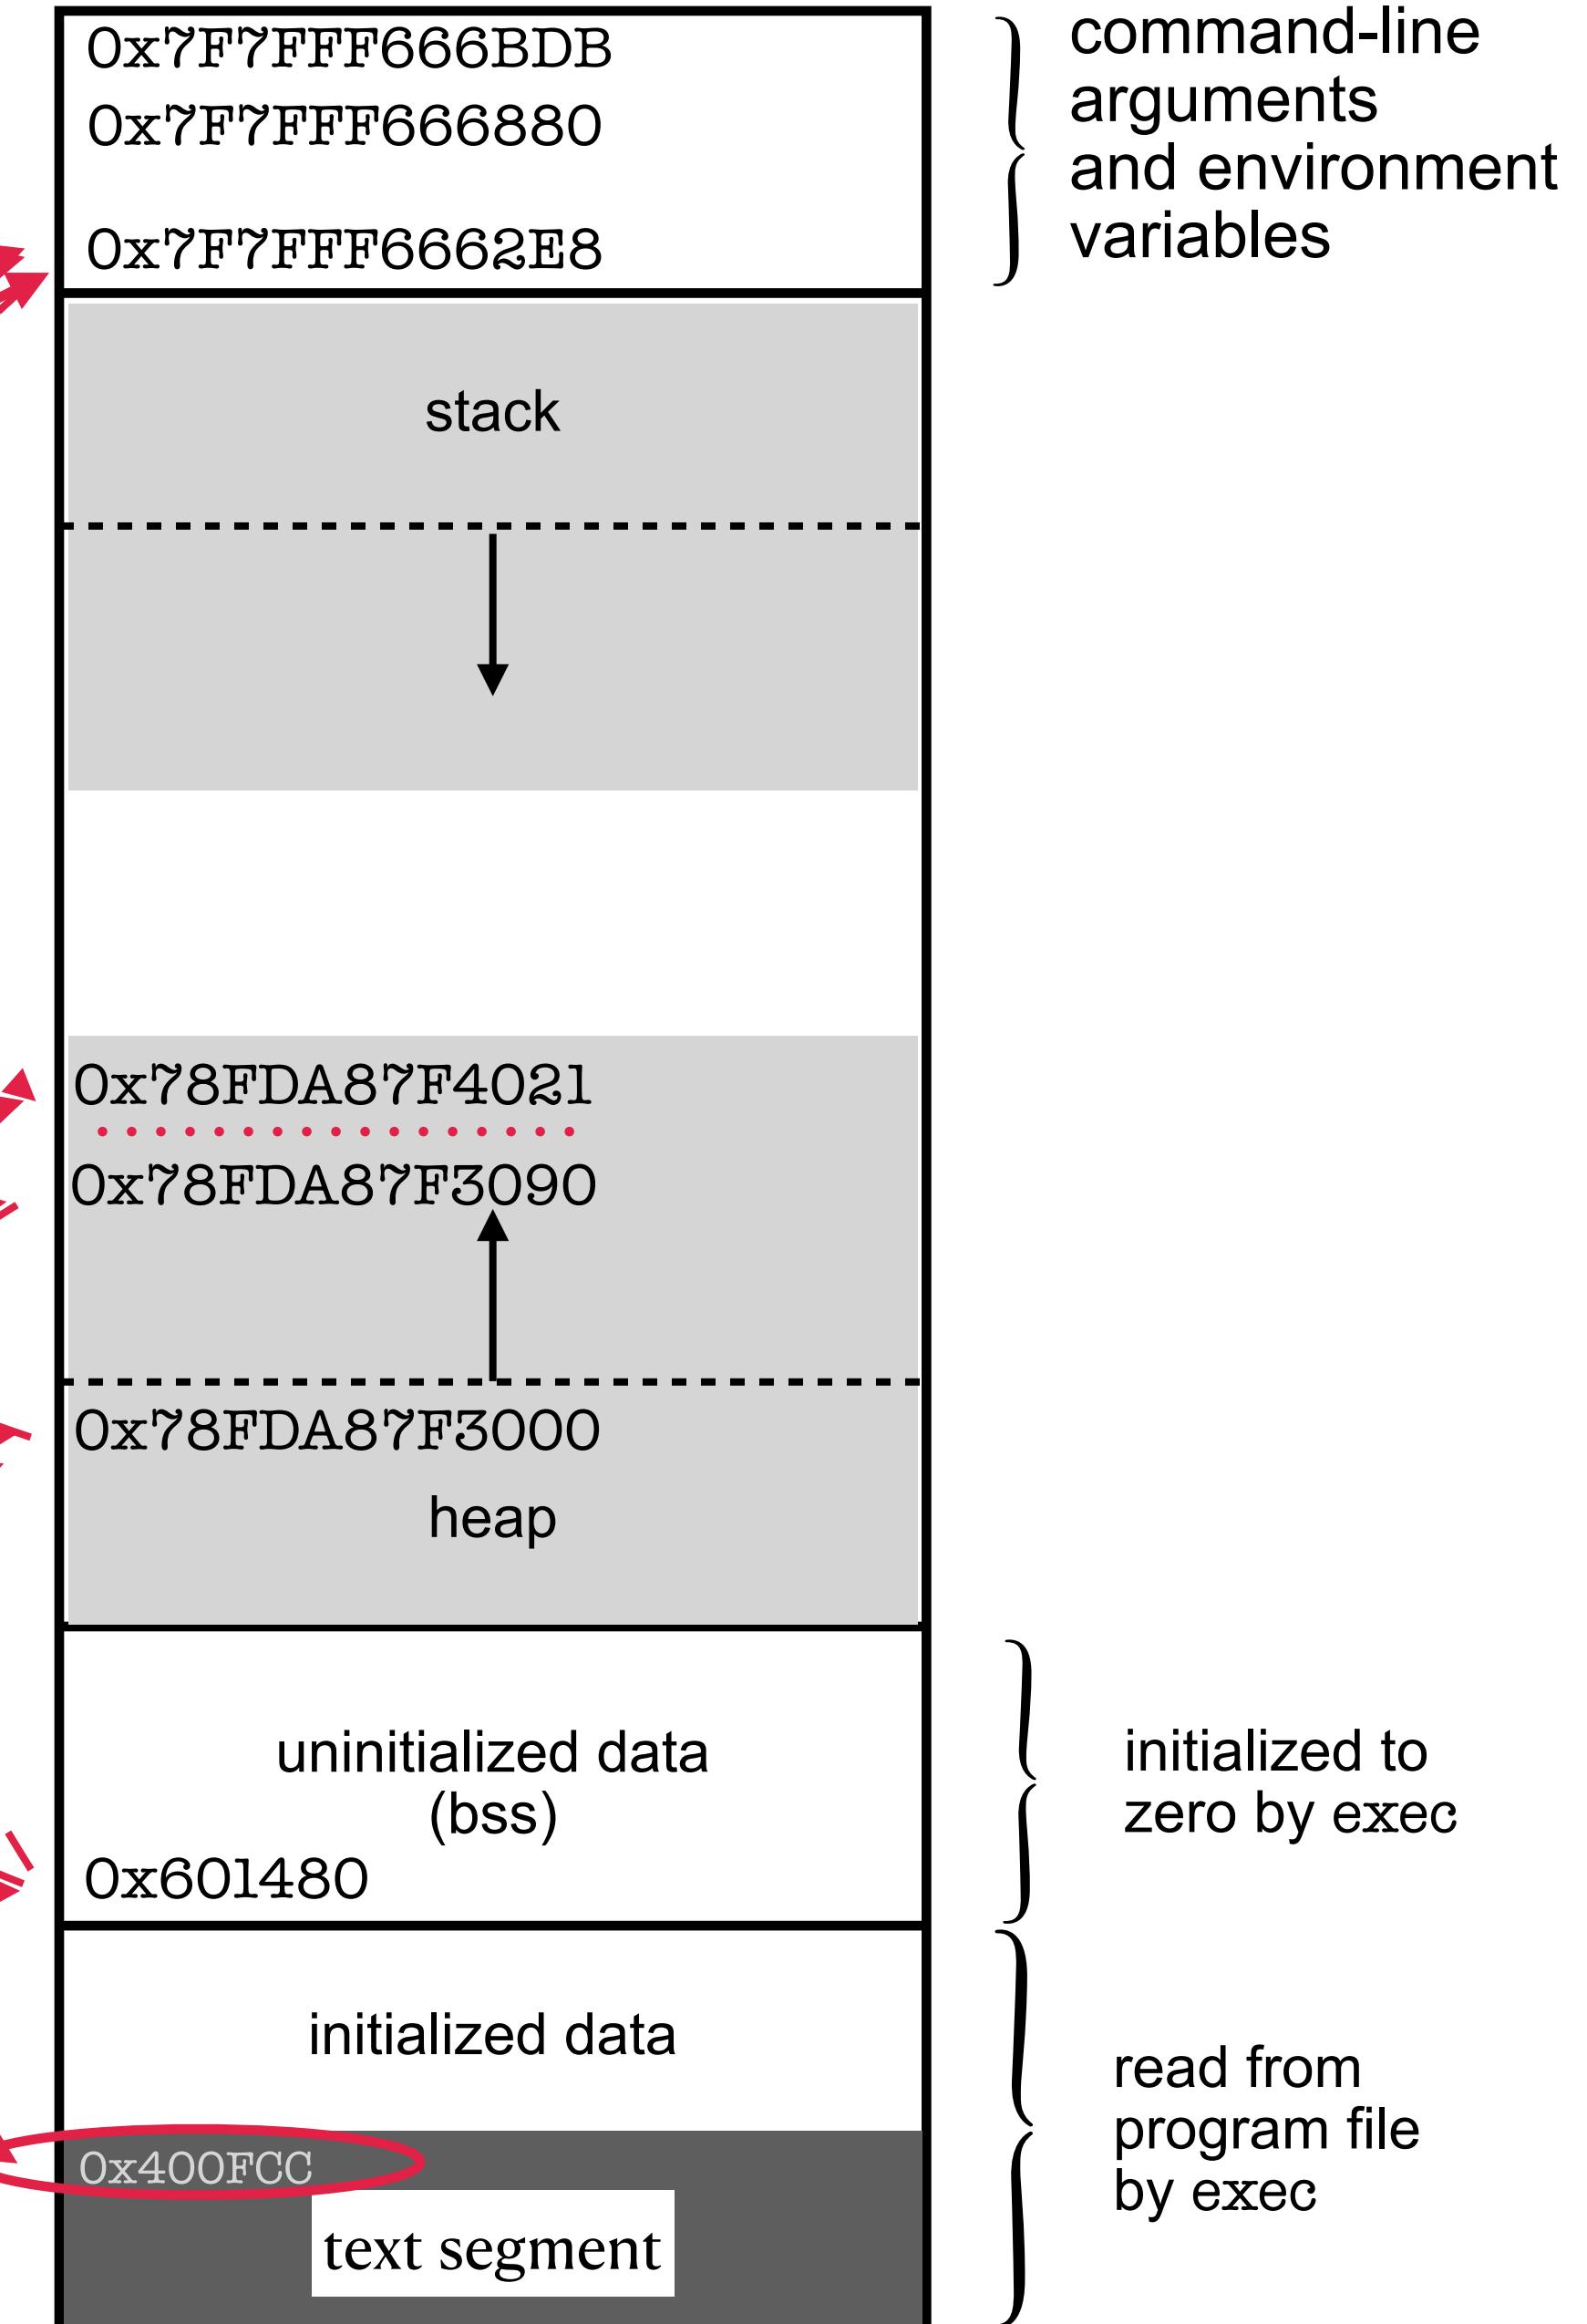
0x7FDA87F3098  environ[19]        (0x400FCC 'BAR=another variable')
0x7FDA87F3090  environ[18]        (0x78FDA87F4021 'FOO=a longer value')
0x7FDA87F3000  environ[0]         (0x78FDA87F4021 'FOO=a longer value')

Uninitialized Data (BSS):
-----
0x  601480  extern char **environ

Heap:
-----
0x7FDA87F2020  malloced area ends
0x7FDA87F2000  malloced area begins
```

high address

low address



Terminal — 80x44

```
[apue$ cc -Wall -Werror -Wextra memory-layout8.c
[apue$ ./a.out
At program start:
0x7FFF8BC928  envp [18]          (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  envp [17]          (0x7FFF8BCE30 'F00=bar')
0x7FFF8BC898  envp [0]           (0x7FFF8BCE30 'F00=bar')

0x7FFF8BC928  environ[18]        (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  environ[17]        (0x7FFF8BCE30 'F00=bar')
0x7FFF8BC898  environ[0]         (0x7FFF8BCE30 'F00=bar')

After setenv(3):
0x7FFF8BC928  envp [18]          (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  envp [17]          (0x7FFF8BCE30 'F00=a longer value')
0x7FFF8BC898  envp [0]           (0x77846A343021 'F00=a longer value')

0x7FFF8BC928  environ[18]        (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  environ[17]        (0x7FFF8BCE30 'F00=a longer value')
0x7FFF8BC898  environ[0]         (0x77846A343021 'F00=a longer value')

After putenv(3):
0x7FFF8BC928  envp [18]          (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  envp [17]          (0x7FFF8BCE30 'F00=a longer value')
0x7FFF8BC898  envp [0]           (0x77846A343021 'F00=a longer value')

0x77846A342098 environ[19]       (0x40106C 'BAR=another variable')
0x77846A342090 environ[18]       (0x77846A343021 'F00=a longer value')
0x77846A342000 environ[0]        (0x77846A343021 'F00=a longer value')

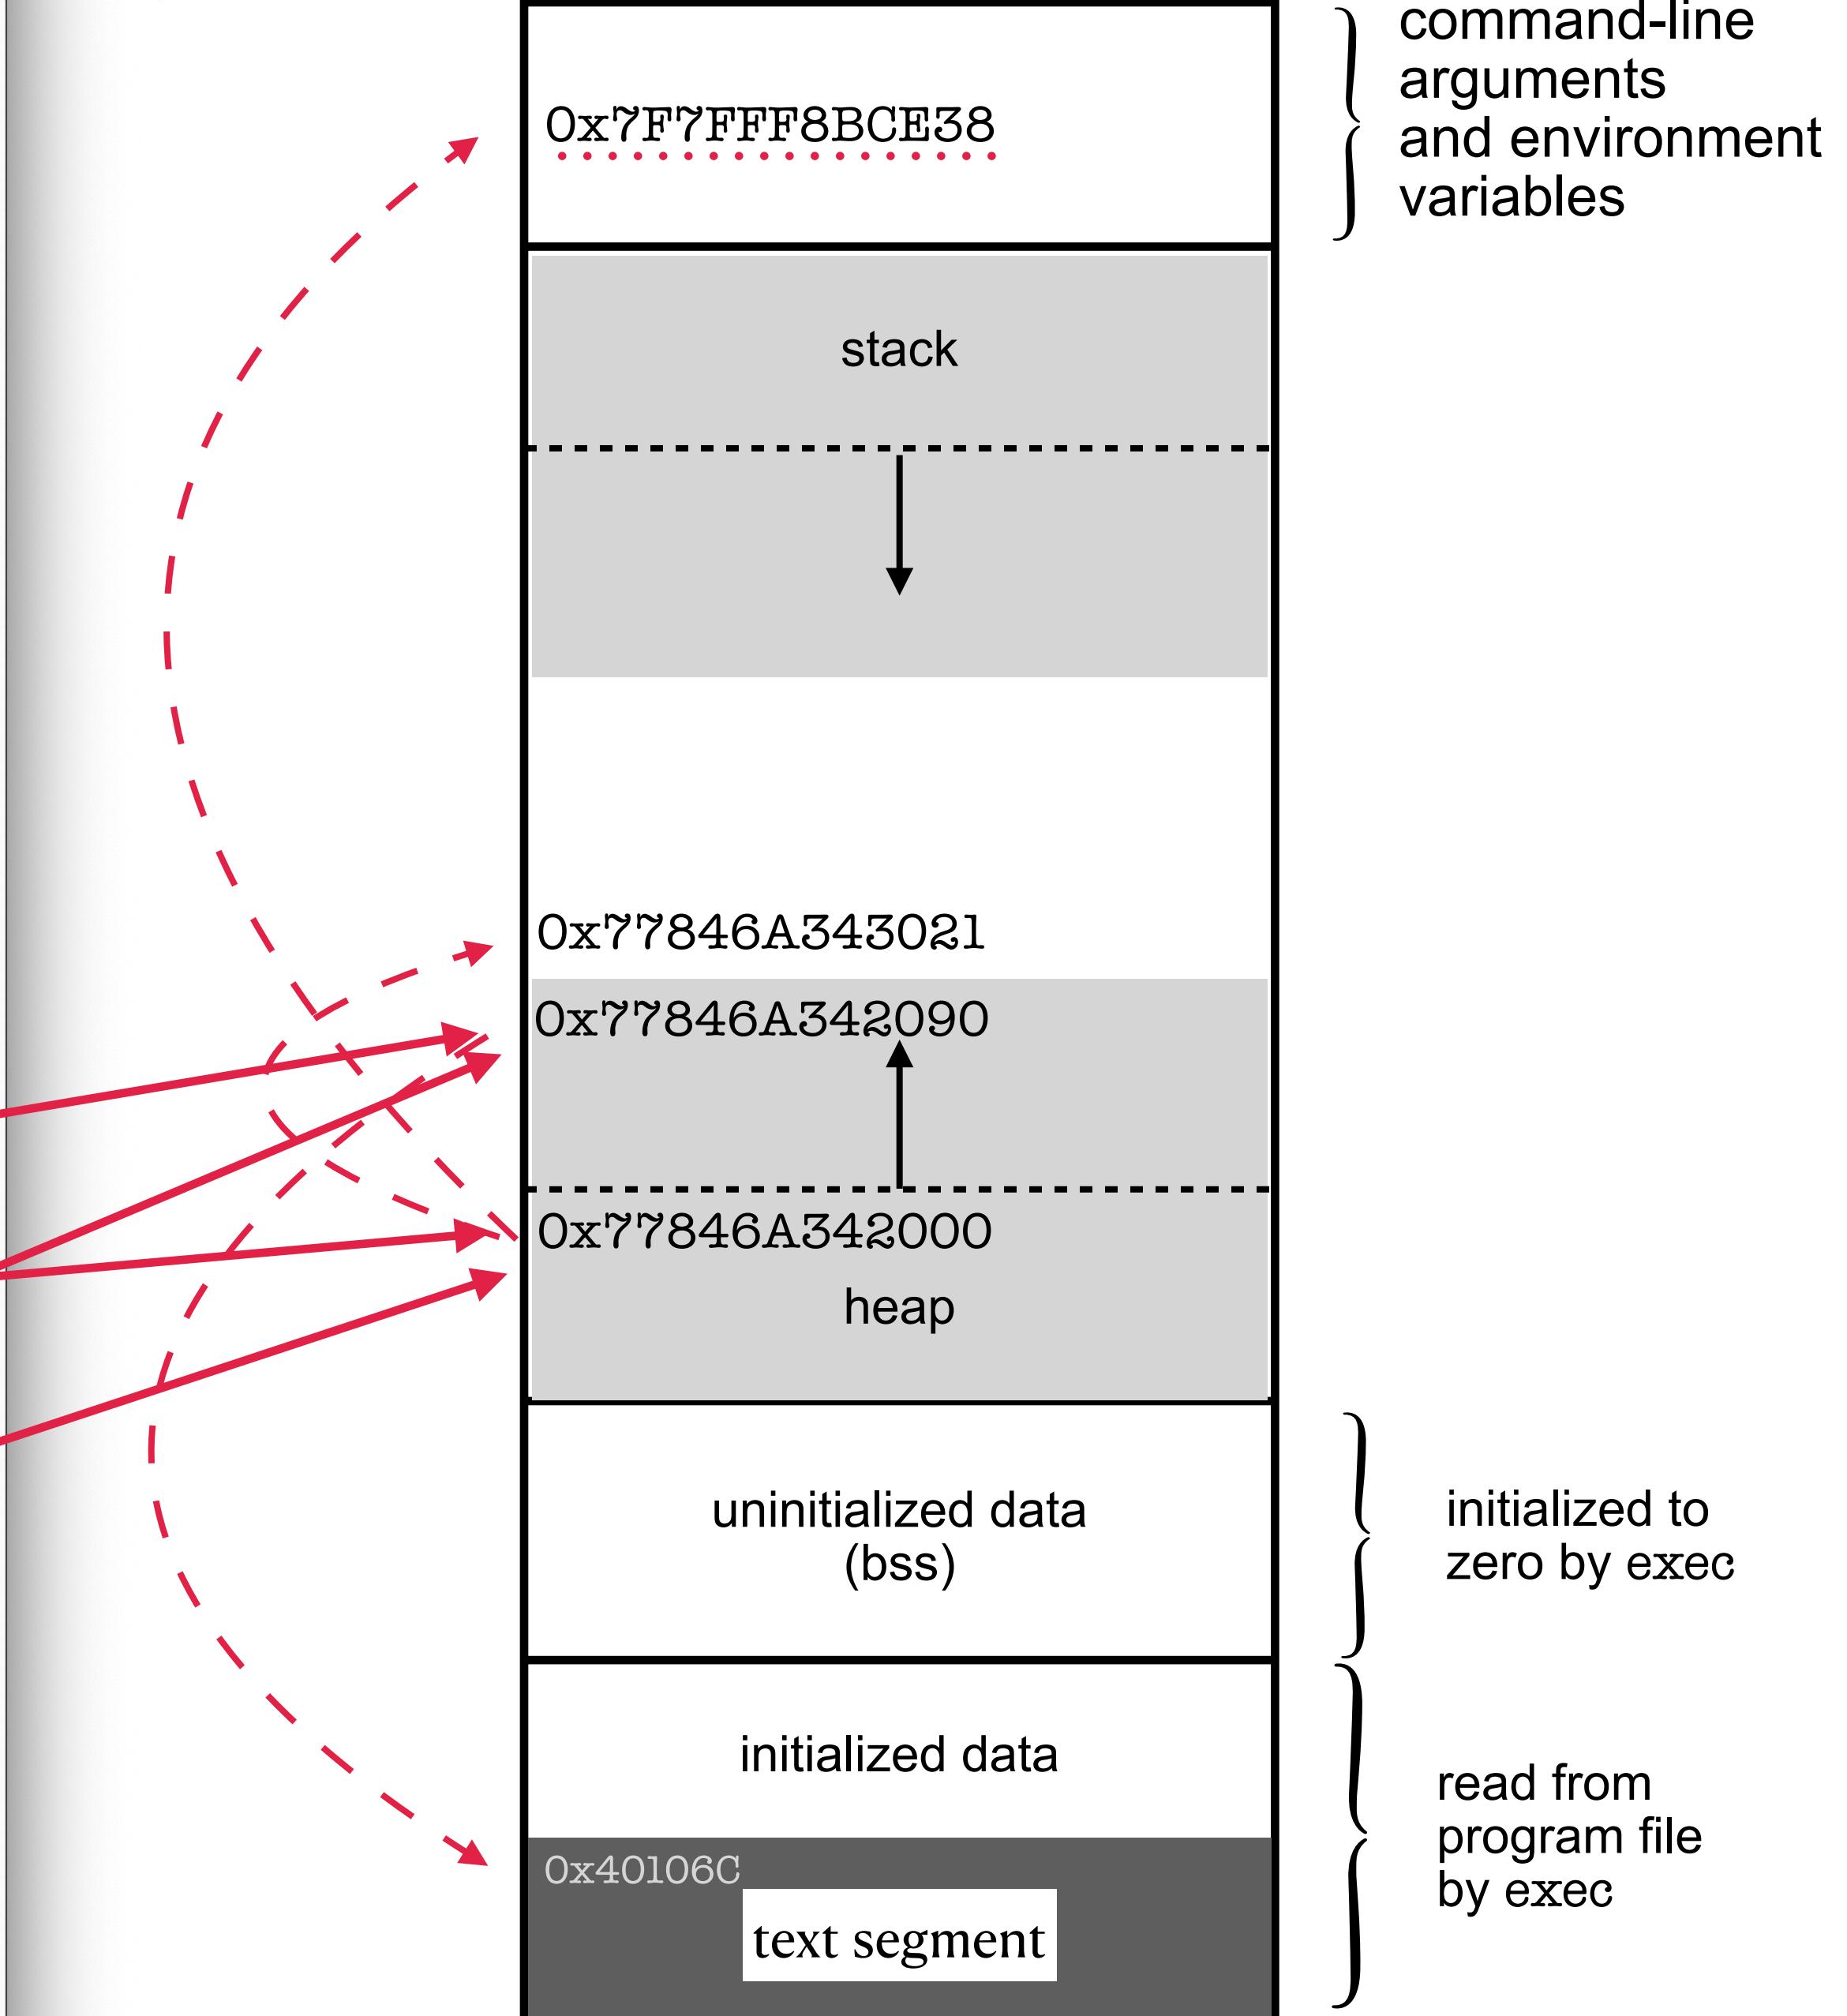
After unsetenv(3):
0x7FFF8BC928  envp [18]          (0x7FFF8BD18B 'SHELL=/bin/sh')
0x7FFF8BC920  envp [17]          (0x7FFF8BCE30 'F00=a longer value')
0x7FFF8BC898  envp [0]           (0x77846A343021 'F00=a longer value')

0x77846A342090 environ[18]       (0x40106C 'BAR=another variable')
0x77846A342088 environ[17]       (0x7FFF8BCE38 'ENV=/home/jschauma/.shrc')
0x77846A342000 environ[0]        (0x7FFF8BCE38 'ENV=/home/jschauma/.shrc')

Stack:
-----
```

high address

low address



The Environment

The process environment consists of an array of strings as KEY=VAR pairs.

The initial environment is set up at the top of the process space and pointed to by the `extern char **environ`.

A third argument `char **envp` may be passed to `main`, which also points to that location.

Setting variables in the environment via `setenv(3)/putenv(3)` may lead to shuffling some or all elements of the `environ` to dynamically allocated memory, but `envp` does *not* get updated!

What programs do with the values in the environment is up to them - but beware, the user can set them to nonsense values!

<https://stevens.netmeister.org/631/env-exercise.html>