

Advanced Programming in the UNIX Environment

Week 07, Segment 4: Signals

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

Signals summarized



Signals

Signals are a way for a process to be notified of *asynchronous* events. Some examples:

- the user hits Ctrl+C (SIGINT)
- the user hits Ctrl+Z (SIGTSTP)
- a background process attempts I/O on the controlling terminal (SIGTTOU / SIGTTIN)
- a timer you set has gone off (SIGALRM)
- a user disconnected from the system (SIGHUP)
- a user resized the terminal window (SIGWINCH)
- ...

See also: `signal(2)/signal(3)/signal(7)` (Note: these man pages vary across platforms!)

Signals

Other ways signals might be generated:

- terminal generated signals (user presses a key combination which causes the terminal driver to generate a signal)
- hardware exceptions (divide by 0, invalid memory references, etc.)
- software conditions (other side of a pipe no longer exists, urgent data has arrived on a network file descriptor, etc.)
- `kill(1)` allows a user to send any signal to any process (if the user is the owner or superuser)
- `kill(2)`, the system call

kill(2)

```
#include <signal.h>  
  
int kill(pid_t pid, int sig);
```

Returns: 0 on success, -1 otherwise

- if $\text{pid} > 0$, then the signal is sent to the process whose PID is pid
- if $\text{pid} == 0$, then the signal is sent to all processes whose process group ID equals the process group ID of the sender
- if $\text{pid} == -1$, then (not POSIX, but BSD and Linux):
 - if $\text{euid} == 0$, the signal is sent to all processes excluding system processes and the process sending the signal
 - else, the signals is sent to all processes with the same uid as the user excluding the process sending the signal
- if $\text{pid} < -1$, then (Linux) the signal is sent to every process in the process group whose ID is $-\text{pid}$
- if $\text{sig} == 0$, no signal is sent, but error checking is performed (*i.e.*, an easy way to test "does pid exist")

Signals

Once we receive a signal, we can do one of several things:

- *Accept the default.* Have the kernel do whatever is defined as the default action for this signal.
- *Ignore it.* (Note: there are some signals which we CANNOT or SHOULD NOT ignore.)
- *Catch it.* That is, have the kernel call a function which we define whenever the signal occurs.
- *Block it.* The delivery of the signal is postponed until it is unblocked.

signal(3)

```
#include <signal.h>  
void (*signal(int sig, void (*func)(int)))(int);
```

Returns: previous signal handler if ok, SIG_ERR otherwise

func can be:

- SIG_IGN, which requests that we ignore the signal signo
- SIG_DFL, which requests that we accept the default action for signal signo
- a pointer to a function to invoke when the signal is received

On some systems/versions, after a signal handler was executed, it was *reset* to SIG_DFL.

signal(3)

```
#include <signal.h>  
typedef void (*sig_t)(int);  
sig_t signal(int sig, sig_t func);
```

Returns: previous signal handler if ok, SIG_ERR otherwise

func can be:

- SIG_IGN, which requests that we ignore the signal signo
- SIG_DFL, which requests that we accept the default action for signal signo
- a pointer to a function to invoke when the signal is received

On some systems/versions, after a signal handler was executed, it was *reset* to SIG_DFL.

signal(3)

```
#include <signal.h>  
void (*signal(int sig, void (*func)(int)))(int);
```

Returns: previous signal handler if ok, SIG_ERR otherwise

func can be:

- SIG_IGN, which requests that we ignore the signal signo
- SIG_DFL, which requests that we accept the default action for signal signo
- a pointer to a function to invoke when the signal is received

On some systems/versions, after a signal handler was executed, it was *reset* to SIG_DFL.

CS631 - Advanced Programming in the UNIX Environment

```
Terminal — 130x24

if (signal(SIGHUP, sig_usr) == SIG_ERR) {
    err(EXIT_FAILURE, "unable to catch SIGHUP");
    /* NOTREACHED */
}

(void)printf("%d\n", getpid());
for ( ; ; ) {
    pause();
    /* Note that the compiler is smart enough to realize we
     * didn't
     * return a value. */
}
apue$ cc -Wall -Werror -Wextra sigusr.c
apue$ ./a.out
1021
Nobody expects SIGUSR1!
A surprising turn of events occurred!
received signal: 1
apue$
```

0 sh

1 sh

sigaction(2)

```
#include <signal.h>
int sigaction(int sig, const struct sigaction * restrict act, struct sigaction * restrict fact);
```

Returns: 0 if ok, -1 otherwise

signal(3) is (nowadays) commonly implemented via sigaction(2):

```
struct sigaction {
    void      (*_sa_handler)(int sig);
    void      (*_sa_sigaction)(int sig, siginfo_t *info, void *ctx);
    sigset_t  sa_mask;
    int       sa_flags;
}
```

More detailed signal concepts

- While executing a signal handler, the signal that triggered it is blocked, but other signals may occur!
- Blocked signals are marked as *pending*; you can inspect the set of pending signals.
- The signal that triggered you entering a signal handler is automatically *unblocked* upon exit from that signal handler, meaning it will be delivered and you will re-enter the same handler right away.
- Multiple signals of the same type arriving in sequence while being blocked may be merged.
- After a `fork(2)`, all signal dispositions and signal masks remain the same in the child as in the parent.
- The `execve(2)` system call reinstates the default action for all signals which were caught, but ignored signals continue to be ignored; the signal mask also remains the same.



```
=> Time for a second interruption.  
^\\In sig_quit, s=2. Now sleeping...  
^\\^\\sig_quit, s=2: exiting  
In sig_quit, s=3. Now sleeping...  
^\\^\\^\\^\\^\\^\\^\\^\\sig_quit, s=3: exiting  
In sig_quit, s=4. Now sleeping...  
sig_quit, s=4: exiting  
Now exiting.  
apue$ ./a.out
```

```
=> Establishing initial signal hander via signal(3).  
^\\In sig_quit, s=1. Now sleeping...  
^CNow in sig_int, s=2. Returning immediately.  
sig_quit, s=2: exiting
```

```
=> Time for a second interruption.  
^\\In sig_quit, s=3. Now sleeping...  
^\\^\\^CNow in sig_int, s=4. Returning immediately.  
sig_quit, s=4: exiting  
In sig_quit, s=5. Now sleeping...  
sig_quit, s=5: exiting  
Now exiting.  
apue$ █
```

```
=> Time for a second interruption.  
^[[1]  Quit (core dumped)      ./a.out  
apue$ ./a.out
```

```
=> Establishing initial signal hander via signal(3).
```

```
=> Establishing a resetting signal hander via signal(3).  
^\\In sig_quit_reset, s=1. Now sleeping...  
^CNow in sig_int, s=2. Returning immediately.  
sig_quit_reset, s=2: exiting
```

```
=> Time for a second interruption.  
^[[1]  Quit (core dumped)      ./a.out  
apue$ ./a.out
```

```
=> Establishing initial signal hander via signal(3).
```

```
=> Establishing a resetting signal hander via signal(3).  
^\\In sig_quit_reset, s=1. Now sleeping...  
^\\^\\^\\^\\^CNow in sig_int, s=2. Returning immediately.  
sig_quit_reset, s=2: exiting  
[1]  Quit (core dumped)      ./a.out  
apue$ █
```



```
=> Unblocking SIGQUIT...
In sig_quit, s=1. Now sleeping...
^\\sig_quit, s=1: exiting
In sig_quit, s=2. Now sleeping...
sig_quit, s=2: exiting
SIGQUIT unblocked - sleeping some more...
^\\In sig_quit, s=3. Now sleeping...
sig_quit, s=3: exiting
Now exiting.
apue$ ./a.out 1
```

```
=> Establishing initial signal hander via signal(3).
```

```
=> Blocking delivery of SIGQUIT...
```

```
=> Now going to sleep for 5 seconds...
```

```
^\\^\\^\\^\\^\\^\\^\\
```

```
=> Checking if any signals are pending...
```

```
=> Unblocking SIGQUIT...
```

```
SIGQUIT unblocked - sleeping some more...
```

```
Now exiting.
```

```
apue$ █
```

Signals

- Signals are notifications of asynchronous events; delivery is similar to a hardware interrupt (current context is saved, a new context is built).
- Signals may be allowed to take the default action (SIG_DFL), be ignored (SIG_IGN), caught (`signal(3)` / `sigaction(2)`), or blocked (`sigprocmask(2)`).
- Arriving signals may be merged, then immediately delivered; different signals may interrupt the current signal handler.
- A number of additional options are available; see `sigaction(2)`.

- If your signal handler can be interrupted, what library functions or system calls are safe to use within a signal handler?