

# **Advanced Programming in the UNIX Environment**

## **Week 09, Segment 5: I/O Multiplexing**

**Department of Computer Science  
Stevens Institute of Technology**

**Jan Schaumann**

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

```
Terminal — 159x39
(j)printf("Client (%s) sent: %s\n", rip, buf);
);
}
}
} while (rval != 0);
(void)close(fd);
}

/*
 * This program uses select() to check that someone is trying to connect
 * before calling accept().
 */
int
main()
{
    int s1, s2;

    s1 = createSocket();
    s2 = createSocket();

    for (;;) {
        handleSocket(s1);
        handleSocket(s2);
    }

    /* NOTREACHED */
}
jschauma@apue$ cc -Wall -Werror -Wextra two-sockets.c
jschauma@apue$ ./a.out
Socket has port #65446
Socket has port #65445
Client (::1) sent: Good Day To You!

Client (::1) sent: What happened to the other connection?

Client (::1) sent: Still blocked, I suppose.

Trying ::1...
Connected to localhost.
Escape character is '^]'.
Good Day To You!
What happened to the other connection?
Still blocked, I suppose.

1 bash
laptop$ telnet apue 65445
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...
Connected to apue.
Escape character is '^]'.
hello there!
Helloooooo!
```

0 bash

1 bash

2 bash

## I/O Multiplexing

---

When handling I/O on multiple file descriptors, we have the following options:

- *blocking mode*: open one fd, block (possibly forever), then test the next fd
- *fork* and use one process for each, communicate using signals or other IPC
- *non-blocking mode*: open one fd, immediately get results, open next fd, immediately get results, sleep for some time
- *asynchronous I/O*: get notified by the kernel when either fd is ready for I/O

## I/O Multiplexing

---

Instead of *blocking forever* (undesirable), using *non-blocking mode* (busy-polling is inefficient) or using *asynchronous I/O* (somewhat limited), we can:

- build a set of file descriptors we're interested in
- call a function that will return if any of the file descriptors are ready for I/O (or a timeout has elapsed)

## select(2)

```
#include <sys/select.h>  
  
int select(int nfds, fd_set * restrict readfds, fd_set * restrict writefds,  
          fd_set * restrict exceptfds, struct timeval * restrict timeout);
```

Returns: # of descriptors if ok, -1 otherwise

select(2) tells us both the total count of descriptors that are ready as well as which ones are ready:

- nfds: which descriptors we're interested in
- readfds, writefds, exceptfds: what conditions we're interested in
- timeout: how long we want to wait
  - tvptr == NULL: wait forever
  - tvptr->tvsec == tvptr->tvusec == 0: don't wait at all

## I/O Multiplexing

---

- file descriptor sets are manipulated using the FD\* functions/macros
- readfds/writefds sets indicate readiness for read/write; exceptfds indicates an exception condition (for example OOB data, certain terminal events)
- EOF means ready for read - read(2) will just return 0 (as usual)
- pselect(2) provides finer-grained timeout control; allows you to specify a signal mask (original signal mask is restored upon return)



Terminal — 159x39

```
+      FD_ZERO(&ready);
+
+      FD_SET(s1, &ready);
+      FD_SET(s2, &ready);
+      timeout.tv_sec = SLEEP;
+      timeout.tv_usec = 0;
+      if (select(s1 > s2 ? s1 + 1 : s2 + 1, &ready, 0, 0, &timeout) <
0) {
+          perror("select");
+          continue;
+      }
+
+      if (FD_ISSET(s1, &ready)) {
+          handleSocket(s1);
+      }
+      if (FD_ISSET(s2, &ready)) {
+          handleSocket(s2);
+      }
+ }

/* NOTREACHED */
}

jschauma@apue$ cc -Wall -Werror -Wextra two-sockets-select.c
jschauma@apue$ ./a.out
Socket has port #65441
Socket has port #65440
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: Hello there!
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: This used to be blocked.
Ending connection
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: This also works.
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: Although that's not very surprising.
```

Ending connection

0 bash

```
* before calling accept().
*/
int
main()
{
    int s1, s2;
    s1 = createSocket();
    s2 = createSocket();
    for (;;) {
        fd_set ready;
        struct timeval timeout;
        FD_ZERO(&ready);
        FD_SET(s1, &ready);
        FD_SET(s2, &ready);
    }
}
```

107,1 2128 77%

```
1 bash
laptop$ telnet apue 65440
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...
Connected to apue.
Escape character is '^]'.
Hello there!
This used to be blocked.
^]
telnet> quit
Connection closed.
laptop$ telnet apue 65441
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...
Connected to apue.
Escape character is '^]'.
This also works.
Although that's not very surprising.
^]
telnet> quit
Connection closed.
laptop$
```

2 bash

```
~  
~  
~  
jschauma@apue$ cc -Wall -Werror -Wextra one-socket-select.c  
jschauma@apue$ ./a.out  
Socket has port #65432  
Idly sitting here, waiting for connections...  
Idly sitting here, waiting for connections...  
Idly sitting here, waiting for connections...  
Client connection from ::1!  
Client (::1) sent: Hello there!  
Client (::1) sent: Note that being connected means the FD set is "ready".  
Client (::1) sent: So we're blocking right now.  
Ending connection from ::1.  
Idly sitting here, waiting for connections...  
Idly sitting here, waiting for connections...  
Idly sitting here, waiting for connections...  
Client connection from ::1!  
Client (::1) sent: But one problem remains.  
Client (::1) sent: Only one connection is being handled.  
Client (::1) sent: We need to disconnect in the first!  
Ending connection from ::1.  
Client connection from 2601:87:4281:1e00:11bb:ea21:91b2:d9c5!  
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: When we're connected multi  
ple times...  
To have our second connection get handled,  
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: Only then do we get to pic  
k up the second connection and process all the data that client sent to us.  
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: I think we need to do bett  
er than that...  
^C  
jschauma@apue$
```

Terminal — 159x39

```
jschauma@apue$ telnet localhost 65432  
Trying ::1...  
Connected to localhost.  
Escape character is '^]'.  
But one problem remains.  
Only one connection is being handled.  
We need to disconnect in the first!  
^]  
telnet> quit  
Connection closed.  
jschauma@apue$
```

1 bash

```
laptop$ telnet apue 65432  
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...  
Connected to apue.  
Escape character is '^]'.  
When we're connected multiple times...  
To have our second connection get handled,  
Only then do we get to pick up the second connection and process all the data t  
hat client sent to us.  
I think we need to do better than that...  
Connection closed by foreign host.  
laptop$
```

2 bash

```
}

void
handleSocket(int s)
{
    int fd;
    pid_t pid;
    struct sockaddr_in6 client;
    socklen_t length;

    length = sizeof(client);
    if ((fd = accept(s, (struct sockaddr *)&client, &length)) < 0) {
        perror("accept");
        return;
    }

jschauma@apue$ cc -Wall -Werror -Wextra one-socket-select-fork.c
jschauma@apue$ ./a.out
Socket has port #65427
Idly sitting here, waiting for connections...
Idly sitting here, waiting for connections...
Client connection from ::1!
Client (::1) sent: Hello from localhost!
Idly sitting here, waiting for connections...
Idly sitting here, waiting for connections...
Idly sitting here, waiting for connections...
Client connection from 2601:87:4281:1e00:11bb:ea21:91b2:d9c5!
Idly sitting here, waiting for connections...
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: Hello from my laptop!
Idly sitting here, waiting for connections...
Idly sitting here, waiting for connections...
Client (::1) sent: Now both clients can simultaneously
Idly sitting here, waiting for connections...
Client (2601:87:4281:1e00:11bb:ea21:91b2:d9c5) sent: communicate with the server without being blocked
Idly sitting here, waiting for connections...
Idly sitting here, waiting for connections...
^C
jschauma@apue$
```

Terminal — 159x39

```
jschauma@apue$ telnet localhost 65432
Trying ::1...
Connected to localhost.
Escape character is '^].
But one problem remains.
Only one connection is being handled.
We need to disconnect in the first!
^]
telnet> quit
Connection closed.
jschauma@apue$ telnet localhost 65427
Trying ::1...
Connected to localhost.
Escape character is '^].
Hello from localhost!
Now both clients can simultaneously
Connection closed by foreign host.
jschauma@apue$
```

1 bash

```
laptop$ telnet apue 65432
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...
Connected to apue.
Escape character is '^].
When we're connected multiple times...
To have our second connection get handled,
Only then do we get to pick up the second connection and process all the data that client sent to us.
I think we need to do better than that...
Connection closed by foreign host.
laptop$ telnet apue 65427
Trying 2601:87:4281:1e00:e12f:43b9:949b:1e0a...
Connected to apue.
Escape character is '^].
Hello from my laptop!
communicate with the server without being blocked
Connection closed by foreign host.
laptop$
```

2 bash

## I/O Multiplexing

---

- use `select(2)` to check multiple file descriptors for I/O readiness
- avoid blocking individual file descriptors by handling I/O in a separate process or thread
- alternative or similar interfaces:
  - `poll(2)`
  - `epoll(2)` (Linux)
  - `kqueue(2)` (\*BSD)
  - <https://libevent.org/> / <http://software.schmorp.de/pkg/libev.html>
- See also: <http://www.kegel.com/c10k.html>